

## 1. Business Problem

### 1.1 Problem Description

Netflix is all about connecting people to the movies they love. To help customers find those movies, they developed world-class movie recommendation system: CinematchSM. Its job is to predict whether someone will enjoy a movie based on how much they liked or disliked other movies. Netflix use those predictions to make personal movie recommendations based on each customer's unique tastes. And while **Cinematch** is doing pretty well, it can always be made better.

Now there are a lot of interesting alternative approaches to how Cinematch works that netflix haven't tried. Some are described in the literature, some aren't. We're curious whether any of these can beat Cinematch by making better predictions. Because, frankly, if there is a much better approach it could make a big difference to our customers and our business.

Credits: <https://www.netflixprize.com/rules.html> (<https://www.netflixprize.com/rules.html>)

### 1.2 Problem Statement

Netflix provided a lot of anonymous rating data, and a prediction accuracy bar that is 10% better than what Cinematch can do on the same training data set. (Accuracy is a measurement of how closely predicted ratings of movies match subsequent actual ratings.)

### 1.3 Sources

- <https://www.netflixprize.com/rules.html> (<https://www.netflixprize.com/rules.html>)
- <https://www.kaggle.com/netflix-inc/netflix-prize-data> (<https://www.kaggle.com/netflix-inc/netflix-prize-data>)
- Netflix blog: <https://medium.com/netflix-techblog/netflix-recommendations-beyond-the-5-stars-part-1-55838468f429> (<https://medium.com/netflix-techblog/netflix-recommendations-beyond-the-5-stars-part-1-55838468f429>) (very nice blog)
- surprise library: <http://surpriselib.com/> (<http://surpriselib.com/>) (we use many models from this library)
- surprise library doc: [http://surprise.readthedocs.io/en/stable/getting\\_started.html](http://surprise.readthedocs.io/en/stable/getting_started.html) ([http://surprise.readthedocs.io/en/stable/getting\\_started.html](http://surprise.readthedocs.io/en/stable/getting_started.html)) (we use many models from this library)
- installing surprise: <https://github.com/NicolasHug/Surprise#installation> (<https://github.com/NicolasHug/Surprise#installation>)
- Research paper: <http://courses.ischool.berkeley.edu/i290-dm/s11/SECURE/a1-koren.pdf> (<http://courses.ischool.berkeley.edu/i290-dm/s11/SECURE/a1-koren.pdf>) (most of our work was inspired by this paper)
- SVD Decomposition : <https://www.youtube.com/watch?v=P5mlg91as1c> (<https://www.youtube.com/watch?v=P5mlg91as1c>)

## 1.4 Real world/Business Objectives and constraints

Objectives:

1. Predict the rating that a user would give to a movie that he has not yet rated.
2. Minimize the difference between predicted and actual rating (RMSE and MAPE)

Constraints:

1. Some form of interpretability.

## 2. Machine Learning Problem

### 2.1 Data

#### 2.1.1 Data Overview

Get the data from : <https://www.kaggle.com/netflix-inc/netflix-prize-data/data> (<https://www.kaggle.com/netflix-inc/netflix-prize-data/data>)

Data files :

- combined\_data\_1.txt
- combined\_data\_2.txt

- combined\_data\_3.txt
- combined\_data\_4.txt
- movie\_titles.csv

The first line of each file [combined\_data\_1.txt, combined\_data\_2.txt, combined\_data\_3.txt, combined\_data\_4.txt] contains the movie id followed by a colon. Each subsequent line in the file corresponds to a rating from a customer and its date in the following format:

CustomerID,Rating,Date

MovieIDs range from 1 to 17770 sequentially.

CustomerIDs range from 1 to 2649429, with gaps. There are 480189 users.

Ratings are on a five star (integral) scale from 1 to 5.

Dates have the format YYYY-MM-DD.

## 2.1.2 Example Data point

1:

```
1488844,3,2005-09-06
822109,5,2005-05-13
885013,4,2005-10-19
30878,4,2005-12-26
823519,3,2004-05-03
893988,3,2005-11-17
124105,4,2004-08-05
1248029,3,2004-04-22
1842128,4,2004-05-09
2238063,3,2005-05-11
1503895,4,2005-05-19
2207774,5,2005-06-06
2590061,3,2004-08-12
2442,3,2004-04-14
543865,4,2004-05-28
1209119,4,2004-03-23
804919,4,2004-06-10
1086807,3,2004-12-28
1711859,4,2005-05-08
372233,5,2005-11-23
1080361,3,2005-03-28
1245640,3,2005-12-19
558634,4,2004-12-14
2165002,4,2004-04-06
1181550,3,2004-02-01
```

1227322,4,2004-02-06  
427928,4,2004-02-26  
814701,5,2005-09-29  
808731,4,2005-10-31  
662870,5,2005-08-24  
337541,5,2005-03-23  
786312,3,2004-11-16  
1133214,4,2004-03-07  
1537427,4,2004-03-29  
1209954,5,2005-05-09  
2381599,3,2005-09-12  
525356,2,2004-07-11  
1910569,4,2004-04-12  
2263586,4,2004-08-20  
2421815,2,2004-02-26  
1009622,1,2005-01-19  
1481961,2,2005-05-24  
401047,4,2005-06-03  
2179073,3,2004-08-29  
1434636,3,2004-05-01  
93986,5,2005-10-06  
1308744,5,2005-10-29  
2647871,4,2005-12-30  
1905581,5,2005-08-16  
2508819,3,2004-05-18  
1578279,1,2005-05-19  
1159695,4,2005-02-15  
2588432,3,2005-03-31  
2423091,3,2005-09-12  
470232,4,2004-04-08  
2148699,2,2004-06-05  
1342007,3,2004-07-16  
466135,4,2004-07-13  
2472440,3,2005-08-13  
1283744,3,2004-04-17  
1927580,4,2004-11-08  
716874,5,2005-05-06  
4326,4,2005-10-29

## 2.2 Mapping the real world problem to a Machine Learning Problem

### 2.2.1 Type of Machine Learning Problem

For a given movie and user we need to predict the rating would be given b

y him/her to the movie.

The given problem is a Recommendation problem

It can also be seen as a Regression problem

## 2.2.2 Performance metric

- Mean Absolute Percentage Error:  
[https://en.wikipedia.org/wiki/Mean\\_absolute\\_percentage\\_error](https://en.wikipedia.org/wiki/Mean_absolute_percentage_error)  
([https://en.wikipedia.org/wiki/Mean\\_absolute\\_percentage\\_error](https://en.wikipedia.org/wiki/Mean_absolute_percentage_error))
- Root Mean Square Error: [https://en.wikipedia.org/wiki/Root-mean-square\\_deviation](https://en.wikipedia.org/wiki/Root-mean-square_deviation)  
([https://en.wikipedia.org/wiki/Root-mean-square\\_deviation](https://en.wikipedia.org/wiki/Root-mean-square_deviation))

## 2.2.3 Machine Learning Objective and Constraints

1. Minimize RMSE.
2. Try to provide some interpretability.

```
In [2]: 1 # this is just to know how much time will it take to run this entire ipython
2 from datetime import datetime
3 # globalstart = datetime.now()
4 import pandas as pd
5 import numpy as np
6 import matplotlib
7 matplotlib.use('nbagg')
8
9 import matplotlib.pyplot as plt
10 plt.rcParams.update({'figure.max_open_warning': 0})
11
12 import seaborn as sns
13 sns.set_style('whitegrid')
14 import os
15 from scipy import sparse
16 from scipy.sparse import csr_matrix
17
18 from sklearn.decomposition import TruncatedSVD
19 from sklearn.metrics.pairwise import cosine_similarity
20 import random
```

# 3. Exploratory Data Analysis

## 3.1 Preprocessing

### 3.1.1 Converting / Merging whole data to required format: u i m i.

... combining, merging more data is required format \_., \_.,  
r\_ij

In [2]:

```

1 start = datetime.now()
2 if not os.path.isfile('data.csv'):
3     # Create a file 'data.csv' before reading it
4     # Read all the files in netflix and store them in one big file('data.csv'
5     # We re reading from each of the four files and appendig each rating to a
6     data = open('data.csv', mode='w')#mode here signifies that we are writing
7
8     row = list() #creating list to store all
9     files=['data_folder/combined_data_1.txt', 'data_folder/combined_data_2.txt
10    'data_folder/combined_data_3.txt', 'data_folder/combined_data_4.tx
11    for file in files:
12        print("Reading ratings from {}".format(file))
13        with open(file) as f:
14            for line in f:
15                del row[:] # you don't have to do this.
16                line = line.strip()
17                if line.endswith(':'):
18                    # All below are ratings for this movie, until another mov
19                    movie_id = line.replace(':', '')
20                else:
21                    row = [x for x in line.split(',')]
22                    row.insert(0, movie_id)
23                    data.write(','.join(row))
24                    data.write('\n')
25            print("Done.\n")
26    data.close()
27    print('Time taken :', datetime.now() - start)

```

Time taken : 0:00:00.004000

In [3]:

```

1 print("creating the dataframe from data.csv file..")
2 df = pd.read_csv('data.csv', sep=',',
3                 names=['movie', 'user', 'rating', 'date'])
4 df.date = pd.to_datetime(df.date)
5 print('Done.\n')
6
7 # we are arranging the ratings according to time.
8 print('Sorting the dataframe by date..')
9 df.sort_values(by='date', inplace=True)
10 print('Done..')

```

creating the dataframe from data.csv file..

Done.

Sorting the dataframe by date..

Done..

```
In [0]: 1 df.head()
```

```
Out[14]:
```

	movie	user	rating	date
56431994	10341	510180	4	1999-11-11
9056171	1798	510180	5	1999-11-11
58698779	10774	510180	3	1999-11-11
48101611	8651	510180	2	1999-11-11
81893208	14660	510180	2	1999-11-11

```
In [0]: 1 df.describe()['rating']
```

```
Out[7]: count    1.004805e+08
mean      3.604290e+00
std       1.085219e+00
min       1.000000e+00
25%      3.000000e+00
50%      4.000000e+00
75%      4.000000e+00
max       5.000000e+00
Name: rating, dtype: float64
```

### 3.1.2 Checking for NaN values

```
In [0]: 1 # just to make sure that all Nan containing rows are deleted..
2 print("No of Nan values in our dataframe : ", sum(df.isnull().any()))
```

No of Nan values in our dataframe : 0

### 3.1.3 Removing Duplicates

```
In [0]: 1 dup_bool = df.duplicated(['movie','user','rating'])
2 dups = sum(dup_bool) # by considering all columns..( including timestamp)
3 print("There are {} duplicate rating entries in the data..".format(dups))
```

There are 0 duplicate rating entries in the data..

### 3.1.4 Basic Statistics (#Ratings, #Users, and #Movies)

```
In [0]: 1 print("Total data ")
2 print("-"*50)
3 print("\nTotal no of ratings :",df.shape[0])
4 print("Total No of Users   :", len(np.unique(df.user)))
5 print("Total No of movies  :", len(np.unique(df.movie)))
```

Total data

-----

Total no of ratings : 100480507  
 Total No of Users : 480189  
 Total No of movies : 17770

## 3.2 Splitting data into Train and Test(80:20)

```
In [0]: 1 if not os.path.isfile('train.csv'):
2     # create the dataframe and store it in the disk for offline purposes..
3     df.iloc[:int(df.shape[0]*0.80)].to_csv("train.csv", index=False)
4
5 if not os.path.isfile('test.csv'):
6     # create the dataframe and store it in the disk for offline purposes..
7     df.iloc[int(df.shape[0]*0.80):].to_csv("test.csv", index=False)
8
9 train_df = pd.read_csv("train.csv", parse_dates=['date'])
10 test_df = pd.read_csv("test.csv")
```

### 3.2.1 Basic Statistics in Train data (#Ratings, #Users, and #Movies)

```
In [0]: 1 # movies = train_df.movie.value_counts()
2 # users = train_df.user.value_counts()
3 print("Training data ")
4 print("-"*50)
5 print("\nTotal no of ratings :",train_df.shape[0])
6 print("Total No of Users   :", len(np.unique(train_df.user)))
7 print("Total No of movies  :", len(np.unique(train_df.movie)))
```

Training data

-----

Total no of ratings : 80384405  
 Total No of Users : 405041  
 Total No of movies : 17424

### 3.2.2 Basic Statistics in Test data (#Ratings, #Users, and #Movies)



```
In [0]: 1 print("Test data ")
2 print("-"*50)
3 print("\nTotal no of ratings :",test_df.shape[0])
4 print("Total No of Users   :", len(np.unique(test_df.user)))
5 print("Total No of movies  :", len(np.unique(test_df.movie)))
```

Test data

-----

Total no of ratings : 20096102  
Total No of Users : 349312  
Total No of movies : 17757

## 3.3 Exploratory Data Analysis on Train data

```
In [0]: 1 # method to make y-axis more readable
2 def human(num, units = 'M'):
3     units = units.lower()
4     num = float(num)
5     if units == 'k':
6         return str(num/10**3) + " K"
7     elif units == 'm':
8         return str(num/10**6) + " M"
9     elif units == 'b':
10        return str(num/10**9) + " B"
```

### 3.3.1 Distribution of ratings

```
In [0]: 1 fig, ax = plt.subplots()
2 plt.title('Distribution of ratings over Training dataset', fontsize=15)
3 sns.countplot(train_df.rating)
4 ax.set_yticklabels([human(item, 'M') for item in ax.get_yticks()])
5 ax.set_ylabel('No. of Ratings(Millions)')
6
7 plt.show()
```

<IPython.core.display.Javascript object>



**Add new column (week day) to the data set for analysis.**

```
In [0]: 1 # It is used to skip the warning ''SettingWithCopyWarning''..  
2 pd.options.mode.chained_assignment = None # default='warn'  
3  
4 train_df['day_of_week'] = train_df.date.dt.weekday_name  
5  
6 train_df.tail()
```


```
Out[17]:
```

	movie	user	rating	date	day_of_week
<b>80384400</b>	12074	2033618	4	2005-08-08	Monday
<b>80384401</b>	862	1797061	3	2005-08-08	Monday
<b>80384402</b>	10986	1498715	5	2005-08-08	Monday
<b>80384403</b>	14861	500016	4	2005-08-08	Monday
<b>80384404</b>	5926	1044015	5	2005-08-08	Monday

### 3.3.2 Number of Ratings per a month

```
In [0]: 1 ax = train_df.resample('m', on='date')['rating'].count().plot()
2 ax.set_title('No of ratings per month (Training data)')
3 plt.xlabel('Month')
4 plt.ylabel('No of ratings(per month)')
5 ax.set_yticklabels([human(item, 'M') for item in ax.get_yticks()])
6 plt.show()
```

<IPython.core.display.Javascript object>



```
In [7]: 1 del df_train
2 del df_test
```

```
-----
NameError                                Traceback (most recent call last)
<ipython-input-7-dee20b8b0016> in <module>()
----> 1 del df_train
      2 del df_test

NameError: name 'df_train' is not defined
```


### 3.3.3 Analysis on the Ratings given by user

```
In [0]: 1 no_of Rated_movies_per_user = train_df.groupby(by='user')['rating'].count().s
        2
        3 no_of Rated_movies_per_user.head()
```

```
Out[20]: user
305344    17112
2439493    15896
387418     15402
1639792     9767
1461435     9447
Name: rating, dtype: int64
```

```
In [0]: 1 fig = plt.figure(figsize=plt.figaspect(.5))
        2
        3 ax1 = plt.subplot(121)
        4 sns.kdeplot(no_of Rated_movies_per_user, shade=True, ax=ax1)
        5 plt.xlabel('No of ratings by user')
        6 plt.title("PDF")
        7
        8 ax2 = plt.subplot(122)
        9 sns.kdeplot(no_of Rated_movies_per_user, shade=True, cumulative=True, ax=ax2)
       10 plt.xlabel('No of ratings by user')
       11 plt.title('CDF')
       12
       13 plt.show()
```

<IPython.core.display.Javascript object>



```
In [0]: 1 no_of Rated movies per user.describe()
```

```
Out[22]: count      405041.000000  
mean         198.459921  
std          290.793238  
min           1.000000  
25%          34.000000  
50%          89.000000  
75%         245.000000  
max        17112.000000  
Name: rating, dtype: float64
```

*There, is something interesting going on with the quantiles..*


```
In [0]: 1 quantiles = no_of Rated movies per user.quantile(np.arange(0,1.01,0.01), inte
```

```

In [0]: 1 plt.title("Quantiles and their Values")
2 quantiles.plot()
3 # quantiles with 0.05 difference
4 plt.scatter(x=quantiles.index[::5], y=quantiles.values[::5], c='orange', label =
5 # quantiles with 0.25 difference
6 plt.scatter(x=quantiles.index[::25], y=quantiles.values[::25], c='m', label =
7 plt.ylabel('No of ratings by user')
8 plt.xlabel('Value at the quantile')
9 plt.legend(loc='best')
10
11 # annotate the 25th, 50th, 75th and 100th percentile values....
12 for x,y in zip(quantiles.index[::25], quantiles[::25]):
13     plt.annotate(s="({} , {})".format(x,y), xy=(x,y), xytext=(x-0.05, y+500)
14                 ,fontweight='bold')
15
16
17 plt.show()

```

<IPython.core.display.Javascript object>



```
In [0]: 1 quantiles[::5]
```

```
Out[25]: 0.00      1
          0.05      7
          0.10     15
          0.15     21
          0.20     27
          0.25     34
          0.30     41
          0.35     50
          0.40     60
          0.45     73
          0.50     89
          0.55    109
          0.60    133
          0.65    163
          0.70    199
          0.75    245
          0.80    307
          0.85    392
          0.90    520
          0.95    749
          1.00   17112
          Name: rating, dtype: int64
```

**how many ratings at the last 5% of all ratings??**

```
In [0]: 1 print('\n No of ratings at last 5 percentile : {}'.format(sum(no_of Rated_Movie
```


```
No of ratings at last 5 percentile : 20305
```

### 3.3.4 Analysis of ratings of a movie given by a user



```
In [0]: 1 no_of_ratings_per_movie = train_df.groupby(by='movie')['rating'].count().sort
        2
        3 fig = plt.figure(figsize=plt.figaspect(.5))
        4 ax = plt.gca()
        5 plt.plot(no_of_ratings_per_movie.values)
        6 plt.title('# RATINGS per Movie')
        7 plt.xlabel('Movie')
        8 plt.ylabel('No of Users who rated a movie')
        9 ax.set_xticklabels([])
       10
       11 plt.show()
```

<IPython.core.display.Javascript object>




- It is very skewed.. just like number of ratings given per user.
  - There are some movies (which are very popular) which are rated by huge number of users.
  - But most of the movies(like 90%) got some hundreds of ratings.

### 3.3.5 Number of ratings on each day of the week


```
In [0]: 1 fig, ax = plt.subplots()
2 sns.countplot(x='day_of_week', data=train_df, ax=ax)
3 plt.title('No of ratings on each day...')
4 plt.ylabel('Total no of ratings')
5 plt.xlabel('')
6 ax.set_yticklabels([human(item, 'M') for item in ax.get_yticks()])
7 plt.show()
```

<IPython.core.display.Javascript object>



```
In [0]: 1 start = datetime.now()
2 fig = plt.figure(figsize=plt.figaspect(.45))
3 sns.boxplot(y='rating', x='day_of_week', data=train_df)
4 plt.show()
5 print(datetime.now() - start)
```

<IPython.core.display.Javascript object>



0:01:10.003761

```
In [0]: 1 avg_week_df = train_df.groupby(by=['day_of_week'])['rating'].mean()
2 print(" Average ratings")
3 print("-"*30)
4 print(avg_week_df)
5 print("\n")
```

Average ratings

```
-----
day_of_week
Friday      3.585274
Monday      3.577250
Saturday    3.591791
Sunday      3.594144
Thursday    3.582463
Tuesday     3.574438
Wednesday   3.583751
Name: rating, dtype: float64
```

### 3.3.6 Creating sparse matrix from data frame

MOVIE_ID	USER_ID	RATING
1	1	3
2	1	4
3	1	2
3	2	1
4	2	4
8	2	2
1	3	3
7	3	1
10	3	5



	1	2	3	4	5	6	7	8	9	10	(movie)
1	3	4	2	-	-	-	-	-	-	-	
2	-	-	1	4	-	-	-	-	-	-	
3	3	-	-	-	-	-	1	-	-	5	
(user)											

#### 3.3.6.1 Creating sparse matrix from train data frame

```

In [0]: 1 start = datetime.now()
2 if os.path.isfile('train_sparse_matrix.npz'):
3     print("It is present in your pwd, getting it from disk...")
4     # just get it from the disk instead of computing it
5     train_sparse_matrix = sparse.load_npz('train_sparse_matrix.npz')
6     print("DONE..")
7 else:
8     print("We are creating sparse_matrix from the dataframe..")
9     # create sparse_matrix and store it for after usage.
10    # csr_matrix(data_values, (row_index, col_index), shape_of_matrix)
11    # It should be in such a way that, MATRIX[row, col] = data
12    train_sparse_matrix = sparse.csr_matrix((train_df.rating.values, (train_d
13                                            train_df.movie.values)),)
14
15    print('Done. It\'s shape is : (user, movie) : ',train_sparse_matrix.shape)
16    print('Saving it into disk for furthur usage..')
17    # save it into disk
18    sparse.save_npz("train_sparse_matrix.npz", train_sparse_matrix)
19    print('Done..\n')
20
21 print(datetime.now() - start)

```

We are creating sparse\_matrix from the dataframe..  
 Done. It's shape is : (user, movie) : (2649430, 17771)  
 Saving it into disk for furthur usage..  
 Done..

0:01:13.804969

### The Sparsity of Train Sparse Matrix

```

In [0]: 1 us,mv = train_sparse_matrix.shape
2 elem = train_sparse_matrix.count_nonzero()
3
4 print("Sparsity Of Train matrix : {} % ".format( (1-(elem/(us*mv))) * 100) )

```

Sparsity Of Train matrix : 99.8292709259195 %

### 3.3.6.2 Creating sparse matrix from test data frame

```

In [0]: 1 start = datetime.now()
2 if os.path.isfile('test_sparse_matrix.npz'):
3     print("It is present in your pwd, getting it from disk...")
4     # just get it from the disk instead of computing it
5     test_sparse_matrix = sparse.load_npz('test_sparse_matrix.npz')
6     print("DONE..")
7 else:
8     print("We are creating sparse_matrix from the dataframe..")
9     # create sparse_matrix and store it for after usage.
10    # csr_matrix(data_values, (row_index, col_index), shape_of_matrix)
11    # It should be in such a way that, MATRIX[row, col] = data
12    test_sparse_matrix = sparse.csr_matrix((test_df.rating.values, (test_df.u
13                                           test_df.movie.values)))
14
15    print('Done. It\'s shape is : (user, movie) : ', test_sparse_matrix.shape)
16    print('Saving it into disk for furthur usage..')
17    # save it into disk
18    sparse.save_npz("test_sparse_matrix.npz", test_sparse_matrix)
19    print('Done..\n')
20
21 print(datetime.now() - start)

```

We are creating sparse\_matrix from the dataframe..  
 Done. It's shape is : (user, movie) : (2649430, 17771)  
 Saving it into disk for furthur usage..  
 Done..

0:00:18.566120

### The Sparsity of Test data Matrix

```

In [0]: 1 us,mv = test_sparse_matrix.shape
2 elem = test_sparse_matrix.count_nonzero()
3
4 print("Sparsity Of Test matrix : {} % ".format( (1-(elem/(us*mv))) * 100) )

```

Sparsity Of Test matrix : 99.95731772988694 %

### 3.3.7 Finding Global average of all movie ratings, Average rating per user, and Average rating per movie

```

In [0]: 1 # get the user averages in dictionary (key: user_id/movie_id, value: avg rating)
        2
        3 def get_average_ratings(sparse_matrix, of_users):
        4
        5     # average ratings of user/axes
        6     ax = 1 if of_users else 0 # 1 - User axes, 0 - Movie axes
        7
        8     # ".A1" is for converting Column_Matrix to 1-D numpy array
        9     sum_of_ratings = sparse_matrix.sum(axis=ax).A1
       10     # Boolean matrix of ratings ( whether a user rated that movie or not)
       11     isRated = sparse_matrix!=0
       12     # no of ratings that each user OR movie..
       13     no_of_ratings = isRated.sum(axis=ax).A1
       14
       15     # max_user and max_movie ids in sparse matrix
       16     u,m = sparse_matrix.shape
       17     # create a dictionary of users and their average ratings..
       18     average_ratings = { i : sum_of_ratings[i]/no_of_ratings[i]
       19                        for i in range(u if of_users else m)
       20                        if no_of_ratings[i] !=0}
       21
       22     # return that dictionary of average ratings
       23     return average_ratings

```

### 3.3.7.1 finding global average of all movie ratings

```

In [0]: 1 train_averages = dict()
        2 # get the global average of ratings in our train set.
        3 train_global_average = train_sparse_matrix.sum()/train_sparse_matrix.count_nonzero()
        4 train_averages['global'] = train_global_average
        5 train_averages

```

Out[36]: {'global': 3.582890686321557}

### 3.3.7.2 finding average rating per user

```

In [0]: 1 train_averages['user'] = get_average_ratings(train_sparse_matrix, of_users=True)
        2 print('\nAverage rating of user 10 : ',train_averages['user'][10])

```

Average rating of user 10 : 3.3781094527363185

### 3.3.7.3 finding average rating per movie

```

In [0]: 1 train_averages['movie'] = get_average_ratings(train_sparse_matrix, of_movies=True)
        2 print('\n Average rating of movie 15 : ',train_averages['movie'][15])


```

Average rating of movie 15 : 3.3038461538461537

### 3.3.7.4 PDF's & CDF's of Avg.Ratings of Users & Movies (In Train Data)

```
In [0]: 1 start = datetime.now()
2 # draw pdfs for average rating per user and average
3 fig, (ax1, ax2) = plt.subplots(nrows=1, ncols=2, figsize=plt.figaspect(.5))
4 fig.suptitle('Avg Ratings per User and per Movie', fontsize=15)
5
6 ax1.set_title('Users-Avg-Ratings')
7 # get the list of average user ratings from the averages dictionary..
8 user_averages = [rat for rat in train_averages['user'].values()]
9 sns.distplot(user_averages, ax=ax1, hist=False,
10               kde_kws=dict(cumulative=True), label='Cdf')
11 sns.distplot(user_averages, ax=ax1, hist=False, label='Pdf')
12
13 ax2.set_title('Movies-Avg-Rating')
14 # get the list of movie average ratings from the dictionary..
15 movie_averages = [rat for rat in train_averages['movie'].values()]
16 sns.distplot(movie_averages, ax=ax2, hist=False,
17               kde_kws=dict(cumulative=True), label='Cdf')
18 sns.distplot(movie_averages, ax=ax2, hist=False, label='Pdf')
19
20 plt.show()
21 print(datetime.now() - start)
```

<IPython.core.display.Javascript object>



0:00:35.003443

## 3.3.8 Cold Start problem

### 3.3.8.1 Cold Start problem with Users



```
In [0]: 1 total_users = len(np.unique(df.user))
2 users_train = len(train_averages['user'])
3 new_users = total_users - users_train
4
5 print('\nTotal number of Users  :', total_users)
6 print('\nNumber of Users in Train data :', users_train)
7 print("\nNo of Users that didn't appear in train data: {}({} %) \n ".format(n
8 np.rc
```

Total number of Users : 480189

Number of Users in Train data : 405041

No of Users that didn't appear in train data: 75148(15.65 %)

We might have to handle **new users ( 75148 )** who didn't appear in train data.

### 3.3.8.2 Cold Start problem with Movies

```
In [0]: 1 total_movies = len(np.unique(df.movie))
2 movies_train = len(train_averages['movie'])
3 new_movies = total_movies - movies_train
4
5 print('\nTotal number of Movies  :', total_movies)
6 print('\nNumber of Users in Train data :', movies_train)
7 print("\nNo of Movies that didn't appear in train data: {}({} %) \n ".format(
8 np.rc
```

Total number of Movies : 17770

Number of Users in Train data : 17424

No of Movies that didn't appear in train data: 346(1.95 %)

We might have to handle **346 movies** (small comparatively) in test data

## 3.4 Computing Similarity matrices

### 3.4.1 Computing User-User Similarity matrix

1. Calculating User User Similarity\_Matrix is **not very easy**(*unless you have huge Computing Power and lots of time*) because of number of. usersbeing lare.
  - You can try if you want to. Your system could crash or the program stops with **Memory Error**

#### 3.4.1.1 Trying with all dimensions (17k dimensions per user)

```

In [0]: 1 from sklearn.metrics.pairwise import cosine_similarity
2
3
4 def compute_user_similarity(sparse_matrix, compute_for_few=False, top = 100,
5                             draw_time_taken=True):
6     no_of_users, _ = sparse_matrix.shape
7     # get the indices of non zero rows(users) from our sparse matrix
8     row_ind, col_ind = sparse_matrix.nonzero()
9     row_ind = sorted(set(row_ind)) # we don't have to
10    time_taken = list() # time taken for finding similar users for an user..
11
12    # we create rows, cols, and data lists.., which can be used to create spa
13    rows, cols, data = list(), list(), list()
14    if verbose: print("Computing top",top,"similarities for each user..")
15
16    start = datetime.now()
17    temp = 0
18
19    for row in row_ind[:top] if compute_for_few else row_ind:
20        temp = temp+1
21        prev = datetime.now()
22
23        # get the similarity row for this user with all other users
24        sim = cosine_similarity(sparse_matrix.getrow(row), sparse_matrix).ravel()
25        # We will get only the top 'top' most similar users and ignore rest
26        top_sim_ind = sim.argsort()[-top:]
27        top_sim_val = sim[top_sim_ind]
28
29        # add them to our rows, cols and data
30        rows.extend([row]*top)
31        cols.extend(top_sim_ind)
32        data.extend(top_sim_val)
33        time_taken.append(datetime.now().timestamp() - prev.timestamp())
34        if verbose:
35            if temp%verb_for_n_rows == 0:
36                print("computing done for {} users [ time elapsed : {} ]".
37                      .format(temp, datetime.now()-start))
38
39
40    # Lets create sparse matrix out of these and return it
41    if verbose: print('Creating Sparse matrix from the computed similarities')
42    #return rows, cols, data
43
44    if draw_time_taken:
45        plt.plot(time_taken, label = 'time taken for each user')
46        plt.plot(np.cumsum(time_taken), label='Total time')
47        plt.legend(loc='best')
48        plt.xlabel('User')
49        plt.ylabel('Time (seconds)')
50        plt.show()
51
52    return sparse.csr_matrix((data, (rows, cols)), shape=(no_of_users, no_of_

```

```
In [0]: 1 start = datetime.now()
2 u_u_sim_sparse, _ = compute_user_similarity(train_sparse_matrix, compute_for_
3                                             verbose=True)
4 print("-"*100)
5 print("Time taken :",datetime.now()-start)
```

Computing top 100 similarities for each user..

computing done for 20 users [ time elapsed : 0:03:20.300488 ]

computing done for 40 users [ time elapsed : 0:06:38.518391 ]


computing done for 60 users [ time elapsed : 0:09:53.143126 ]

computing done for 80 users [ time elapsed : 0:13:10.080447 ]

computing done for 100 users [ time elapsed : 0:16:24.711032 ]

Creating Sparse matrix from the computed similarities

<IPython.core.display.Javascript object>



Time taken : 0:16:33.618931

### 3.4.1.2 Trying with reduced dimensions (Using TruncatedSVD for dimensionality reduction of user vector)

- We have **405,041 users** in our training set and computing similarities between them..( **17K dimensional vector..**) is time consuming..
- From above plot, It took roughly **8.88 sec** for computing similar users for **one user**
- We have **405,041 users** with us in training set.
- $405041 \times 8.88 = 3596764.08 \text{ sec} = 59946.068 \text{ min} = 999.101133333 \text{ hours} = 41.62921$ 
  - Even if we run on 4 cores parallelly (a typical system now a days), It will still take almost **10 and 1/2 days**.

IDEA: Instead, we will try to reduce the dimensions using SVD, so that **it might** speed up the process...

```
In [0]: 1 from datetime import datetime
2 from sklearn.decomposition import TruncatedSVD
3
4 start = datetime.now()
5
6 # initialize the algorithm with some parameters..
7 # All of them are default except n_components. n_iter is for Randomized SVD so
8 netflix_svd = TruncatedSVD(n_components=500, algorithm='randomized', random_s
9 trunc_svd = netflix_svd.fit_transform(train_sparse_matrix)
10
11 print(datetime.now()-start)
```

0:29:07.069783

Here,

- $\Sigma \leftarrow (\text{netflix\_svd.singular\_values\_})$
- $V^T \leftarrow (\text{netflix\_svd.components\_})$
- $U$  is not returned. instead **Projection\_of\_X** onto the new vectorspace is returned.
- It uses **randomized svd** internally, which returns **All 3 of them separately**. Use that instead..


```
In [0]: 1 expl_var = np.cumsum(netflix_svd.explained_variance_ratio_)
```

```

In [0]: 1 fig, (ax1, ax2) = plt.subplots(nrows=1, ncols=2, figsize=plt.figaspect(.5))
2
3 ax1.set_ylabel("Variance Explained", fontsize=15)
4 ax1.set_xlabel("# Latent Facors", fontsize=15)
5 ax1.plot(expl_var)
6 # annotate some (latentfactors, expl_var) to make it clear
7 ind = [1, 2,4,8,20, 60, 100, 200, 300, 400, 500]
8 ax1.scatter(x = [i-1 for i in ind], y = expl_var[[i-1 for i in ind]], c='#ff3
9 for i in ind:
10     ax1.annotate(s="({}, {})".format(i, np.round(expl_var[i-1], 2)), xy=(i-
11         xytext = ( i+20, expl_var[i-1] - 0.01), fontweight='bold')
12
13 change_in_expl_var = [expl_var[i+1] - expl_var[i] for i in range(len(expl_var
14 ax2.plot(change_in_expl_var)
15
16
17
18 ax2.set_ylabel("Gain in Var_Expl with One Additional LF", fontsize=10)
19 ax2.yaxis.set_label_position("right")
20 ax2.set_xlabel("# Latent Facors", fontsize=20)
21
22 plt.show()

```

<IPython.core.display.Javascript object>



```
In [0]: 1 for i in ind:
        2     print("{} {}".format(i, np.round(expl_var[i-1], 2)))
```

(1, 0.23)  
 (2, 0.26)  
 (4, 0.3)  
 (8, 0.34)  
 (20, 0.38)  
 (60, 0.44)  
 (100, 0.47)  
 (200, 0.53)  
 (300, 0.57)  
 (400, 0.61)  
 (500, 0.64)

I think 500 dimensions is good enough

- By just taking **(20 to 30)** latent factors, explained variance that we could get is **20 %**.
- To take it to **60%**, we have to take **almost 400 latent factors**. It is not fare.
- It basically is the **gain of variance explained**, if we **add one additional latent factor to it**.
- By adding one by one latent factor too it, the **\_gain in explained variance** with that addition is decreasing. (Obviously, because they are sorted that way).
- **LHS Graph:**
  - **x** --- ( No of latent factors ),
  - **y** --- ( The variance explained by taking x latent factors)
- **More decrease in the line (RHS graph) :**
  - We are getting more explained variance than before.
- **Less decrease in that line (RHS graph) :**
  - We are not getting benefitted from adding latent factor further. This is what is shown in the plots.
- **RHS Graph:**
  - **x** --- ( No of latent factors ),
  - **y** --- ( Gain in Expl\_Var by taking one additional latent factor)

```
In [0]: 1 # Let's project our Original U_M matrix into into 500 Dimensional space...
        2 start = datetime.now()
        3 trunc_matrix = train_sparse_matrix.dot(netflix_svd.components_.T)
        4 print(datetime.now()- start)
```

0:00:45.670265

```
In [0]: 1 type(trunc_matrix), trunc_matrix.shape
```

Out[53]: (numpy.ndarray, (2649430, 500))

- Let's convert this to actual sparse matrix and store it for future purposes

```
In [0]: 1 if not os.path.isfile('trunc_sparse_matrix.npz'):
2         # create that sparse matrix
3         trunc_sparse_matrix = sparse.csr_matrix(trunc_matrix)
4         # Save this truncated sparse matrix for later usage..
5         sparse.save_npz('trunc_sparse_matrix', trunc_sparse_matrix)
6     else:
7         trunc_sparse_matrix = sparse.load_npz('trunc_sparse_matrix.npz')
```

```
In [0]: 1 trunc_sparse_matrix.shape
```

```
Out[55]: (2649430, 500)
```



```
In [0]: 1 start = datetime.now()
2 trunc_u_u_sim_matrix, _ = compute_user_similarity(trunc_sparse_matrix, comput
3                                                    verb_for_n_rows=10)
4 print("-"*50)
5 print("time:",datetime.now()-start)
```

Computing top 50 similarities for each user..

computing done for 10 users [ time elapsed : 0:02:09.746324 ]

computing done for 20 users [ time elapsed : 0:04:16.017768 ]


computing done for 30 users [ time elapsed : 0:06:20.861163 ]

computing done for 40 users [ time elapsed : 0:08:24.933316 ]

computing done for 50 users [ time elapsed : 0:10:28.861485 ]

Creating Sparse matrix from the computed similarities

<IPython.core.display.Javascript object>



-----  
time: 0:10:52.658092

**: This is taking more time for each user than Original one.**

- from above plot, It took almost **12.18** for computing similar users for **one user**
- We have **405041 users** with us in training set.
- $405041 \times 12.18 == 4933399.38 \text{ sec} == 82223.323 \text{ min} == 1370.388716667 \text{ h}$

- Even we run on 4 cores parallelly (a typical system now a days), It will still take almost **(14 - 15)** days.



- **Why did this happen...??**

- Just think about it. It's not that difficult.

-----(*sparse & dense.....get it ??*)-----

### Is there any other way to compute user user similarity..??

-An alternative is to compute similar users for a particular user, whenever required (**ie., Run time**)

- We maintain a binary Vector for users, which tells us whether we already computed or not..
- **\*\*\*If not\*\*\* :**
  - Compute top (let's just say, 1000) most similar users for this given user, and add this to our datastructure, so that we can just access it (similar users) without recomputing it again.
  -
- **\*\*\*If It is already Computed\*\*\*:**
  - Just get it directly from our datastructure, which has that information.
  - In production time, We might have to recompute similarities, if it is computed a long time ago. Because user preferences changes over time. If we could maintain some kind of Timer, which when expires, we have to update it ( recompute it ).
  -
- **\*\*\*Which datastructure to use:\*\*\***
  - It is purely implementation dependant.
  - One simple method is to maintain a **\*\*Dictionary Of Dictionaries\*\***.
    - 
    - **\*\*key : \_userid\_**
    - **\_\_value\_\_ : \_Again a dictionary\_**
      - **\_\_key\_\_ : \_Similar User\_**
      - **\_\_value\_\_ : \_Similarity Value\_**

### 3.4.2 Computing Movie-Movie Similarity matrix

```
In [0]: 1 start = datetime.now()
2 if not os.path.isfile('m_m_sim_sparse.npz'):
3     print("It seems you don't have that file. Computing movie_movie similarity...")
4     start = datetime.now()
5     m_m_sim_sparse = cosine_similarity(X=train_sparse_matrix.T, dense_output=True)
6     print("Done..")
7     # store this sparse matrix in disk before using it. For future purposes.
8     print("Saving it to disk without the need of re-computing it again.. ")
9     sparse.save_npz("m_m_sim_sparse.npz", m_m_sim_sparse)
10    print("Done..")
11 else:
12     print("It is there, We will get it.")
13     m_m_sim_sparse = sparse.load_npz("m_m_sim_sparse.npz")
14     print("Done ...")
15
16 print("It's a ",m_m_sim_sparse.shape," dimensional matrix")
17
18 print(datetime.now() - start)
```

It seems you don't have that file. Computing movie\_movie similarity...  
Done..  
Saving it to disk without the need of re-computing it again..  
Done..  
It's a (17771, 17771) dimensional matrix  
0:10:02.736054

```
In [0]: 1 m_m_sim_sparse.shape
```

Out[59]: (17771, 17771)

- Even though we have similarity measure of each movie, with all other movies, We generally don't care much about least similar movies.
- Most of the times, only top\_xxx similar items matters. It may be 10 or 100.
- We take only those top similar movie ratings and store them in a saperate dictionary.

```
In [0]: 1 movie_ids = np.unique(m_m_sim_sparse.nonzero()[1])
```

```

In [0]: 1 start = datetime.now()
        2 similar_movies = dict()
        3 for movie in movie_ids:
        4     # get the top similar movies and store them in the dictionary
        5     sim_movies = m_m_sim_sparse[movie].toarray().ravel().argsort()[::-1][1:]
        6     similar_movies[movie] = sim_movies[:100]
        7 print(datetime.now() - start)
        8
        9 # just testing similar movies for movie_15
       10 similar_movies[15]

```

0:00:33.411700

```

Out[62]: array([ 8279,  8013, 16528,  5927, 13105, 12049,  4424, 10193, 17590,
                4549,  3755,   590, 14059, 15144, 15054,  9584,  9071,  6349,
               16402,  3973,  1720,  5370, 16309,  9376,  6116,  4706,  2818,
                778, 15331,  1416, 12979, 17139, 17710,  5452,  2534,   164,
               15188,  8323,  2450, 16331,  9566, 15301, 13213, 14308, 15984,
               10597,  6426,  5500,  7068,  7328,  5720,  9802,   376, 13013,
                8003, 10199,  3338, 15390,  9688, 16455, 11730,  4513,   598,
               12762,  2187,   509,  5865,  9166, 17115, 16334,  1942,  7282,
               17584,  4376,  8988,  8873,  5921,  2716, 14679, 11947, 11981,
                4649,   565, 12954, 10788, 10220, 10963,  9427,  1690,  5107,
                7859,  5969,  1510,  2429,   847,  7845,  6410, 13931,  9840,
                3706])

```

### 3.4.3 Finding most similar movies using similarity matrix

**Does Similarity really works as the way we expected...?**

*Let's pick some random movie and check for its similar movies....*

```
In [0]: 1 # First Let's Load the movie details into soe dataframe..
2 # movie details are in 'netflix/movie_titles.csv'
3
4 movie_titles = pd.read_csv("data_folder/movie_titles.csv", sep=',', header =
5                             names=['movie_id', 'year_of_release', 'title'], ve
6                             index_col = 'movie_id', encoding = "ISO-8859-1")
7
8 movie_titles.head()
```

Tokenization took: 4.50 ms  
 Type conversion took: 165.72 ms  
 Parser memory cleanup took: 0.01 ms

```
Out[64]:
```

	year_of_release	title
movie_id		
1	2003.0	Dinosaur Planet
2	2004.0	Isle of Man TT 2004 Review
3	1997.0	Character
4	1994.0	Paula Abdul's Get Up & Dance
5	2004.0	The Rise and Fall of ECW

### Similar Movies for 'Vampire Journals'

```
In [0]: 1 mv_id = 67
2
3 print("\nMovie ----->",movie_titles.loc[mv_id].values[1])
4
5 print("\nIt has {} Ratings from users.".format(train_sparse_matrix[:,mv_id].g
6
7 print("\nWe have {} movies which are similarto this  and we will get only top
```

Movie -----> Vampire Journals


It has 270 Ratings from users.

We have 17284 movies which are similarto this and we will get only top most..

```
In [0]: 1 similarities = m_m_sim_sparse[mv_id].toarray().ravel()
2
3 similar_indices = similarities.argsort()[::-1][1:]
4
5 similarities[similar_indices]
6
7 sim_indices = similarities.argsort()[::-1][1:] # It will sort and reverse the
8                                                # and return its indices(movie
```

```
In [0]: 1 plt.plot(similarities[sim_indices], label='All the ratings')
2         plt.plot(similarities[sim_indices[:100]], label='top 100 similar movies')
3         plt.title("Similar Movies of {}".format(mv_id), fontsize=20)
4         plt.xlabel("Movies (Not Movie_Ids)", fontsize=15)
5         plt.ylabel("Cosine Similarity", fontsize=15)
6         plt.legend()
7         plt.show()
```

<IPython.core.display.Javascript object>



**Top 10 similar movies**

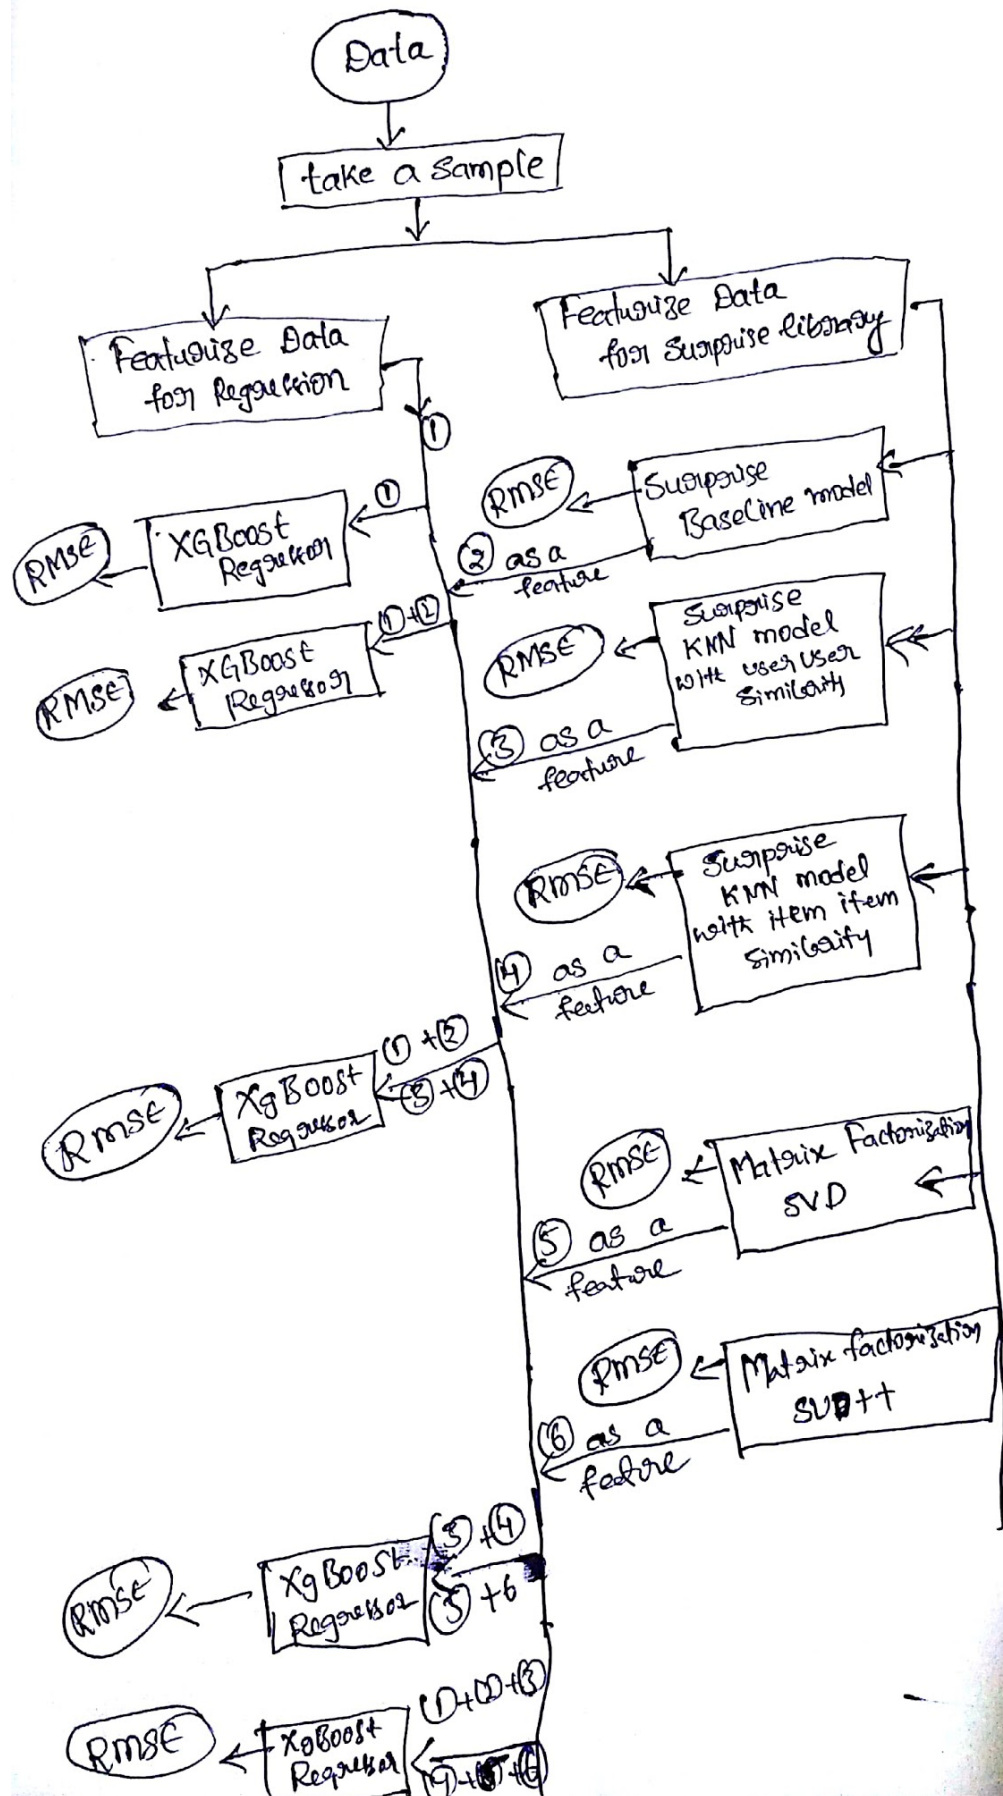
```
In [0]: 1 movie_titles.loc[sim_indices[:10]]
```

```
Out[68]:
```

	year_of_release	title
movie_id		
323	1999.0	Modern Vampires
4044	1998.0	Subspecies 4: Bloodstorm
1688	1993.0	To Sleep With a Vampire
13962	2001.0	Dracula: The Dark Prince
12053	1993.0	Dracula Rising
16279	2002.0	Vampires: Los Muertos
4667	1996.0	Vampirella
1900	1997.0	Club Vampire
13873	2001.0	The Breed
15867	2003.0	Dracula II: Ascension

Similarly, we can *find similar users* and compare how similar they are.

## 4. Machine Learning Models





```

In [0]: 1 def get_sample_sparse_matrix(sparse_matrix, no_users, no_movies, path, verbose
2       """
3           It will get it from the 'path' if it is present or It will create
4           and store the sampled sparse matrix in the path specified.
5       """
6
7       # get (row, col) and (rating) tuple from sparse_matrix...
8       row_ind, col_ind, ratings = sparse.find(sparse_matrix)
9       users = np.unique(row_ind)
10      movies = np.unique(col_ind)
11
12      print("Original Matrix : (users, movies) -- ({} {})".format(len(users), len(movies)))
13      print("Original Matrix : Ratings -- {}".format(len(ratings)))
14
15      # It just to make sure to get same sample everytime we run this program..
16      # and pick without replacement....
17      np.random.seed(15)
18      sample_users = np.random.choice(users, no_users, replace=False)
19      sample_movies = np.random.choice(movies, no_movies, replace=False)
20      # get the boolean mask or these sampled_items in original row/col_inds..
21      mask = np.logical_and( np.isin(row_ind, sample_users),
22                             np.isin(col_ind, sample_movies) )
23
24      sample_sparse_matrix = sparse.csr_matrix((ratings[mask], (row_ind[mask],
25                                                         col_ind[mask]),
26                                                         shape=(max(sample_users)+1, max(sample_movies)+1)))
27
28      if verbose:
29          print("Sampled Matrix : (users, movies) -- ({} {})".format(len(sample_users), len(sample_movies)))
30          print("Sampled Matrix : Ratings --", format(ratings[mask].shape[0]))
31
32      print('Saving it into disk for further usage..')
33      # save it into disk
34      sparse.save_npz(path, sample_sparse_matrix)
35      if verbose:
36          print('Done..\n')
37
38      return sample_sparse_matrix

```

## 4.1 Sampling Data

### 4.1.1 Build sample train data from the train data

```
In [0]: 1 start = datetime.now()
2 path = "sample/small/sample_train_sparse_matrix.npz"
3 if os.path.isfile(path):
4     print("It is present in your pwd, getting it from disk....")
5     # just get it from the disk instead of computing it
6     sample_train_sparse_matrix = sparse.load_npz(path)
7     print("DONE..")
8 else:
9     # get 10k users and 1k movies from available data
10    sample_train_sparse_matrix = get_sample_sparse_matrix(train_sparse_matrix,
11                                                         path = path)
12
13 print(datetime.now() - start)
```

It is present in your pwd, getting it from disk....

DONE..

0:00:00.035179

### 4.1.2 Build sample test data from the test data

```
In [0]: 1 start = datetime.now()
2
3 path = "sample/small/sample_test_sparse_matrix.npz"
4 if os.path.isfile(path):
5     print("It is present in your pwd, getting it from disk....")
6     # just get it from the disk instead of computing it
7     sample_test_sparse_matrix = sparse.load_npz(path)
8     print("DONE..")
9 else:
10    # get 5k users and 500 movies from available data
11    sample_test_sparse_matrix = get_sample_sparse_matrix(test_sparse_matrix,
12                                                         path = "sample/small/sample_
13 print(datetime.now() - start)
```

It is present in your pwd, getting it from disk....

DONE..

0:00:00.028740

## 4.2 Finding Global Average of all movie ratings, Average rating per User, and Average rating per Movie (from sampled train)

```
In [0]: 1 sample_train_averages = dict()
```

### 4.2.1 Finding Global Average of all movie ratings

```
In [0]: 1 # get the global average of ratings in our train set.  
2 global_average = sample_train_sparse_matrix.sum()/sample_train_sparse_matrix.  
3 sample_train_averages['global'] = global_average  
4 sample_train_averages
```

```
Out[13]: {'global': 3.581679377504138}
```

```
In [1]: 1 from xgboost import XGBRegressor  
2
```

## 4.2.2 Finding Average rating per User

```
In [0]: 1 sample_train_averages['user'] = get_average_ratings(sample_train_sparse_matri  
2 print('\nAverage rating of user 1515220 :',sample_train_averages['user'][1515
```

Average rating of user 1515220 : 3.9655172413793105

## 4.2.3 Finding Average rating per Movie

```
In [0]: 1 sample_train_averages['movie'] = get_average_ratings(sample_train_sparse_mat  
2 print('\n AVerage rating of movie 15153 :',sample_train_averages['movie'][151
```

Average rating of movie 15153 : 2.6458333333333335

## 4.3 Featurizing data

```
In [0]: 1 print('\n No of ratings in Our Sampled train matrix is : {}'.format(sample  
2 print('\n No of ratings in Our Sampled test matrix is : {}'.format(sample_
```

No of ratings in Our Sampled train matrix is : 129286

No of ratings in Our Sampled test matrix is : 7333

### 4.3.1 Featurizing data for regression problem

#### 4.3.1.1 Featurizing train data

```
In [0]: 1 # get users, movies and ratings from our samples train sparse matrix  
        2 sample_train_users, sample_train_movies, sample_train_ratings = sparse.find(s
```

```

In [0]: 1 #####
2 # It took me almost 10 hours to prepare this train dataset.#
3 #####
4 start = datetime.now()
5 if os.path.isfile('sample/small/reg_train.csv'):
6     print("File already exists you don't have to prepare again..." )
7 else:
8     print('preparing {} tuples for the dataset..\\n'.format(len(sample_train_r
9     with open('sample/small/reg_train.csv', mode='w') as reg_data_file:
10         count = 0
11         for (user, movie, rating) in zip(sample_train_users, sample_train_mo
12             st = datetime.now()
13             #     print(user, movie)
14             #----- Ratings of "movie" by similar users of "us
15             # compute the similar Users of the "user"
16             user_sim = cosine_similarity(sample_train_sparse_matrix[user], sa
17             top_sim_users = user_sim.argsort()[::-1][1:] # we are ignoring 'T
18             # get the ratings of most similar users for this movie
19             top_ratings = sample_train_sparse_matrix[top_sim_users, movie].to
20             # we will make it's length "5" by adding movie averages to .
21             top_sim_users_ratings = list(top_ratings[top_ratings != 0][:5])
22             top_sim_users_ratings.extend([sample_train_averages['movie'][movi
23             #     print(top_sim_users_ratings, end=" ")
24
25
26             #----- Ratings by "user" to similar movies of "m
27             # compute the similar movies of the "movie"
28             movie_sim = cosine_similarity(sample_train_sparse_matrix[:,movie]
29             top_sim_movies = movie_sim.argsort()[::-1][1:] # we are ignoring
30             # get the ratings of most similar movie rated by this user..
31             top_ratings = sample_train_sparse_matrix[user, top_sim_movies].to
32             # we will make it's length "5" by adding user averages to.
33             top_sim_movies_ratings = list(top_ratings[top_ratings != 0][:5])
34             top_sim_movies_ratings.extend([sample_train_averages['user'][user
35             #     print(top_sim_movies_ratings, end=" : -- ")
36
37             #-----prepare the row to be stores in a file-----
38             row = list()
39             row.append(user)
40             row.append(movie)
41             # Now add the other features to this data...
42             row.append(sample_train_averages['global']) # first feature
43             # next 5 features are similar_users "movie" ratings
44             row.extend(top_sim_users_ratings)
45             # next 5 features are "user" ratings for similar_movies
46             row.extend(top_sim_movies_ratings)
47             # Avg_user rating
48             row.append(sample_train_averages['user'][user])
49             # Avg_movie rating
50             row.append(sample_train_averages['movie'][movie])
51
52             # finalley, The actual Rating of this user-movie pair...
53             row.append(rating)
54             count = count + 1
55
56             # add rows to the file opened..

```

```

57         reg_data_file.write(','.join(map(str, row)))
58         reg_data_file.write('\n')
59         if (count)%10000 == 0:
60             # print(','.join(map(str, row)))
61             print("Done for {} rows----- {}".format(count, datetime.now()))
62
63
64     print(datetime.now() - start)

```

preparing 129286 tuples for the dataset..

```

Done for 10000 rows----- 0:53:13.974716
Done for 20000 rows----- 1:47:58.228942
Done for 30000 rows----- 2:42:46.963119
Done for 40000 rows----- 3:36:44.807894
Done for 50000 rows----- 4:28:55.311500
Done for 60000 rows----- 5:24:18.493104
Done for 70000 rows----- 6:17:39.669922
Done for 80000 rows----- 7:11:23.970879
Done for 90000 rows----- 8:05:33.787770
Done for 100000 rows----- 9:00:25.463562
Done for 110000 rows----- 9:51:28.530010
Done for 120000 rows----- 10:42:05.382141
11:30:13.699183

```

### Reading from the file to make a Train\_dataframe

```

In [0]: 1 reg_train = pd.read_csv('sample/small/reg_train.csv', names = ['user', 'movie', 'rating', 'timestamp'])
        2 reg_train.head()

```

```

Out[19]:

```

	user	movie	GAvg	sur1	sur2	sur3	sur4	sur5	smr1	smr2	smr3	smr4	smr5	UA
0	53406	33	3.581679	4.0	5.0	5.0	4.0	1.0	5.0	2.0	5.0	3.0	1.0	3.3703
1	99540	33	3.581679	5.0	5.0	5.0	4.0	5.0	3.0	4.0	4.0	3.0	5.0	3.5555
2	99865	33	3.581679	5.0	5.0	4.0	5.0	3.0	5.0	4.0	4.0	5.0	4.0	3.7142
3	101620	33	3.581679	2.0	3.0	5.0	5.0	4.0	4.0	3.0	3.0	4.0	5.0	3.5844
4	112974	33	3.581679	5.0	5.0	5.0	5.0	5.0	3.0	5.0	5.0	5.0	3.0	3.7500

- **GAvg** : Average rating of all the ratings
- **Similar users rating of this movie:**
  - sur1, sur2, sur3, sur4, sur5 ( top 5 similar users who rated that movie.. )
- **Similar movies rated by this user:**
  - smr1, smr2, smr3, smr4, smr5 ( top 5 similar movies rated by this movie.. )
- **UAvg** : User's Average rating
- **MAvg** : Average rating of this movie

- **rating** : Rating of this movie by this user.
- 

#### 4.3.1.2 Featurizing test data

```
In [0]: 1 # get users, movies and ratings from the Sampled Test  
        2 sample_test_users, sample_test_movies, sample_test_ratings = sparse.find(samp
```

```
In [0]: 1 sample_train_averages['global']
```

```
Out[21]: 3.581679377504138
```

```

In [0]: 1 start = datetime.now()
2
3 if os.path.isfile('sample/small/reg_test.csv'):
4     print("It is already created...")
5 else:
6
7     print('preparing {} tuples for the dataset..\\n'.format(len(sample_test_ra
8 with open('sample/small/reg_test.csv', mode='w') as reg_data_file:
9     count = 0
10    for (user, movie, rating) in zip(sample_test_users, sample_test_movi
11        st = datetime.now()
12
13        #----- Ratings of "movie" by similar users of "user"
14        #print(user, movie)
15        try:
16            # compute the similar Users of the "user"
17            user_sim = cosine_similarity(sample_train_sparse_matrix[user]
18            top_sim_users = user_sim.argsort()[::-1][1:] # we are ignorin
19            # get the ratings of most similar users for this movie
20            top_ratings = sample_train_sparse_matrix[top_sim_users, movie
21            # we will make it's length "5" by adding movie averages to .
22            top_sim_users_ratings = list(top_ratings[top_ratings != 0][:5
23            top_sim_users_ratings.extend([sample_train_averages['movie']][
24            # print(top_sim_users_ratings, end="--")
25
26        except (IndexError, KeyError):
27            # It is a new User or new Movie or there are no ratings for g
28            ##### Cold Start Problem #####
29            top_sim_users_ratings.extend([sample_train_averages['global']
30            #print(top_sim_users_ratings)
31        except:
32            print(user, movie)
33            # we just want KeyErrors to be resolved. Not every Exception.
34            raise
35
36
37
38        #----- Ratings by "user" to similar movies of "m
39        try:
40            # compute the similar movies of the "movie"
41            movie_sim = cosine_similarity(sample_train_sparse_matrix[:,mo
42            top_sim_movies = movie_sim.argsort()[::-1][1:] # we are ignor
43            # get the ratings of most similar movie rated by this user..
44            top_ratings = sample_train_sparse_matrix[user, top_sim_movies
45            # we will make it's length "5" by adding user averages to.
46            top_sim_movies_ratings = list(top_ratings[top_ratings != 0][:
47            top_sim_movies_ratings.extend([sample_train_averages['user']][
48            #print(top_sim_movies_ratings)
49        except (IndexError, KeyError):
50            #print(top_sim_movies_ratings, end=" : -- ")
51            top_sim_movies_ratings.extend([sample_train_averages['global']
52            #print(top_sim_movies_ratings)
53        except :
54            raise
55
56        #-----prepare the row to be stores in a file-----

```



```

57     row = list()
58     # add usser and movie name first
59     row.append(user)
60     row.append(movie)
61     row.append(sample_train_averages['global']) # first feature
62     #print(row)
63     # next 5 features are similar_users "movie" ratings
64     row.extend(top_sim_users_ratings)
65     #print(row)
66     # next 5 features are "user" ratings for similar_movies
67     row.extend(top_sim_movies_ratings)
68     #print(row)
69     # Avg_user rating
70     try:
71         row.append(sample_train_averages['user'][user])
72     except KeyError:
73         row.append(sample_train_averages['global'])
74     except:
75         raise
76     #print(row)
77     # Avg_movie rating
78     try:
79         row.append(sample_train_averages['movie'][movie])
80     except KeyError:
81         row.append(sample_train_averages['global'])
82     except:
83         raise
84     #print(row)
85     # finalley, The actual Rating of this user-movie pair...
86     row.append(rating)
87     #print(row)
88     count = count + 1
89
90     # add rows to the file opened..
91     reg_data_file.write(','.join(map(str, row)))
92     #print(','.join(map(str, row)))
93     reg_data_file.write('\n')
94     if (count)%1000 == 0:
95         #print(','.join(map(str, row)))
96         print("Done for {} rows----- {}".format(count, datetime.now()))
97     print("",datetime.now() - start)

```

preparing 7333 tuples for the dataset..

```

Done for 1000 rows----- 0:04:29.293783
Done for 2000 rows----- 0:08:57.208002
Done for 3000 rows----- 0:13:30.333223
Done for 4000 rows----- 0:18:04.050813
Done for 5000 rows----- 0:22:38.671673
Done for 6000 rows----- 0:27:09.697009
Done for 7000 rows----- 0:31:41.933568
0:33:12.529731

```

**Reading from the file to make a test dataframe**

```
In [0]: 1 reg_test_df = pd.read_csv('sample/small/reg_test.csv', names = ['user', 'movie',
2                                     'rating', 'sur1', 'sur2', 'sur3', 'sur4', 'sur5', 'smr1', 'smr2', 'smr3', 'smr4', 'smr5',
3                                     'MAvg', 'UAvg', 'GAvg'])
4 reg_test_df.head(4)
```

```
Out[30]:
```

	user	movie	GAvg	sur1	sur2	sur3	sur4	sur5	smr1	smr2
0	808635	71	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679
1	941866	71	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679
2	1737912	71	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679
3	1849204	71	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679

- **GAvg** : Average rating of all the ratings
- **Similar users rating of this movie:**
  - sur1, sur2, sur3, sur4, sur5 ( top 5 simiular users who rated that movie.. )
- **Similar movies rated by this user:**
  - smr1, smr2, smr3, smr4, smr5 ( top 5 simiular movies rated by this movie.. )
- **UAvg** : User AVerage rating
- **MAvg** : Average rating of this movie
- **rating** : Rating of this movie by this user.

## 4.3.2 Transforming data for Surprise models

```
In [0]: 1 from surprise import Reader, Dataset
```

### 4.3.2.1 Transforming train data

- We can't give raw data (movie, user, rating) to train the model in Surprise library.
- They have a saperate format for TRAIN and TEST data, which will be useful for training the models like SVD, KNNBaseLineOnly....etc.,in Surprise.
- We can form the trainset from a file, or from a Pandas DataFrame.  
[http://surprise.readthedocs.io/en/stable/getting\\_started.html#load-dom-dataframe-py](http://surprise.readthedocs.io/en/stable/getting_started.html#load-dom-dataframe-py)  
[\(http://surprise.readthedocs.io/en/stable/getting\\_started.html#load-dom-dataframe-py\)](http://surprise.readthedocs.io/en/stable/getting_started.html#load-dom-dataframe-py)

```
In [0]: 1 # It is to specify how to read the dataframe.
        2 # for our dataframe, we don't have to specify anything extra..
        3 reader = Reader(rating_scale=(1,5))
        4
        5 # create the traindata from the dataframe...
        6 train_data = Dataset.load_from_df(reg_train[['user', 'movie', 'rating']], reader)
        7
        8 # build the trainset from traindata.. It is of dataset format from surprise
        9 trainset = train_data.build_full_trainset()
```

#### 4.3.2.2 Transforming test data

- Testset is just a list of (user, movie, rating) tuples. (Order in the tuple is important)

```
In [0]: 1 testset = list(zip(reg_test_df.user.values, reg_test_df.movie.values, reg_test_df.rating.values))
        2 testset[:3]
```

```
Out[35]: [(808635, 71, 5), (941866, 71, 4), (1737912, 71, 3)]
```

## 4.4 Applying Machine Learning models

- Global dictionary that stores rmse and mape for all the models....
  - It stores the metrics in a dictionary of dictionaries

**keys** : model names(string)

**value**: dict(**key** : metric, **value** : value )

```
In [0]: 1 models_evaluation_train = dict()
        2 models_evaluation_test = dict()
        3
        4 models_evaluation_train, models_evaluation_test
```

```
Out[36]: ({}, {})
```

Utility functions for running regression models

```

In [0]: 1 # to get rmse and mape given actual and predicted ratings..
2 def get_error_metrics(y_true, y_pred):
3     rmse = np.sqrt(np.mean([ (y_true[i] - y_pred[i])**2 for i in range(len(y_
4     mape = np.mean(np.abs( (y_true - y_pred)/y_true )) * 100
5     return rmse, mape
6
7 #####
8 #####
9 def run_xgboost(algo, x_train, y_train, x_test, y_test, verbose=True):
10     """
11     It will return train_results and test_results
12     """
13
14     # dictionaries for storing train and test results
15     train_results = dict()
16     test_results = dict()
17
18
19     # fit the model
20     print('Training the model..')
21     start = datetime.now()
22     algo.fit(x_train, y_train, eval_metric = 'rmse')
23     print('Done. Time taken : {}'.format(datetime.now()-start))
24     print('Done \n')
25
26     # from the trained model, get the predictions....
27     print('Evaluating the model with TRAIN data...')
28     start = datetime.now()
29     y_train_pred = algo.predict(x_train)
30     # get the rmse and mape of train data...
31     rmse_train, mape_train = get_error_metrics(y_train.values, y_train_pred)
32
33     # store the results in train_results dictionary..
34     train_results = {'rmse': rmse_train,
35                     'mape' : mape_train,
36                     'predictions' : y_train_pred}
37
38     #####
39     # get the test data predictions and compute rmse and mape
40     print('Evaluating Test data')
41     y_test_pred = algo.predict(x_test)
42     rmse_test, mape_test = get_error_metrics(y_true=y_test.values, y_pred=y_t
43     # store them in our test results dictionary.
44     test_results = {'rmse': rmse_test,
45                    'mape' : mape_test,
46                    'predictions':y_test_pred}
47
48     if verbose:
49         print('\nTEST DATA')
50         print('-'*30)
51         print('RMSE : ', rmse_test)
52         print('MAPE : ', mape_test)
53
54     # return these train and test results...
55     return train_results, test_results

```

## Utility functions for Surprise modes

```

In [0]: 1 # it is just to makesure that all of our algorithms should produce same result
2 # everytime they run...
3
4 my_seed = 15
5 random.seed(my_seed)
6 np.random.seed(my_seed)
7
8 #####
9 # get (actual_list , predicted_list) ratings given list
10 # of predictions (prediction is a class in Surprise).
11 #####
12 def get_ratings(predictions):
13     actual = np.array([pred.r_ui for pred in predictions])
14     pred = np.array([pred.est for pred in predictions])
15
16     return actual, pred
17
18 #####
19 # get 'rmse' and 'mape' , given list of prediction objects
20 #####
21 def get_errors(predictions, print_them=False):
22
23     actual, pred = get_ratings(predictions)
24     rmse = np.sqrt(np.mean((pred - actual)**2))
25     mape = np.mean(np.abs(pred - actual)/actual)
26
27     return rmse, mape*100
28
29 #####
30 # It will return predicted ratings, rmse and mape of both train and test data
31 #####
32 def run_surprise(algo, trainset, testset, verbose=True):
33     '''
34         return train_dict, test_dict
35
36         It returns two dictionaries, one for train and the other is for test
37         Each of them have 3 key-value pairs, which specify 'rmse', 'mape'
38     '''
39     start = datetime.now()
40     # dictionaries that stores metrics for train and test..
41     train = dict()
42     test = dict()
43
44     # train the algorithm with the trainset
45     st = datetime.now()
46     print('Training the model...')
47     algo.fit(trainset)
48     print('Done. time taken : {} \n'.format(datetime.now()-st))
49
50     # ----- Evaluating train data-----#
51     st = datetime.now()
52     print('Evaluating the model with train data..')
53     # get the train predictions (list of prediction class inside Surprise)
54     train_preds = algo.test(trainset.build_testset())
55     # get predicted ratings from the train predictions..
56     train_actual_ratings, train_pred_ratings = get_ratings(train_preds)

```

```

57 # get 'rmse' and 'mape' from the train predictions.
58 train_rmse, train_mape = get_errors(train_preds)
59 print('time taken : {}'.format(datetime.now()-st))
60
61 if verbose:
62     print('-'*15)
63     print('Train Data')
64     print('-'*15)
65     print("RMSE : {}\nMAPE : {}".format(train_rmse, train_mape))
66
67 #store them in the train dictionary
68 if verbose:
69     print('adding train results in the dictionary..')
70 train['rmse'] = train_rmse
71 train['mape'] = train_mape
72 train['predictions'] = train_pred_ratings
73
74 #----- Evaluating Test data-----#
75 st = datetime.now()
76 print('\nEvaluating for test data...')
77 # get the predictions( list of prediction classes) of test data
78 test_preds = algo.test(testset)
79 # get the predicted ratings from the list of predictions
80 test_actual_ratings, test_pred_ratings = get_ratings(test_preds)
81 # get error metrics from the predicted and actual ratings
82 test_rmse, test_mape = get_errors(test_preds)
83 print('time taken : {}'.format(datetime.now()-st))
84
85 if verbose:
86     print('-'*15)
87     print('Test Data')
88     print('-'*15)
89     print("RMSE : {}\nMAPE : {}".format(test_rmse, test_mape))
90 # store them in test dictionary
91 if verbose:
92     print('storing the test results in test dictionary...')
93 test['rmse'] = test_rmse
94 test['mape'] = test_mape
95 test['predictions'] = test_pred_ratings
96
97 print('\n'+ '-'*45)
98 print('Total time taken to run this algorithm :', datetime.now() - start)
99
100 # return two dictionaries train and test
101 return train, test

```

#### 4.4.1 XGBoost with initial 13 features

```
In [0]: 1 import xgboost as xgb
```

```
In [0]: 1 # prepare Train data
2 x_train = reg_train.drop(['user', 'movie', 'rating'], axis=1)
3 y_train = reg_train['rating']
4
5 # Prepare Test data
6 x_test = reg_test_df.drop(['user', 'movie', 'rating'], axis=1)
7 y_test = reg_test_df['rating']
8
9 # initialize Our first XGBoost model...
10 first_xgb = xgb.XGBRegressor(silent=False, n_jobs=13, random_state=15, n_estimators=1000)
11 train_results, test_results = run_xgboost(first_xgb, x_train, y_train, x_test, y_test)
12
13 # store the results in models_evaluations dictionaries
14 models_evaluation_train['first_algo'] = train_results
15 models_evaluation_test['first_algo'] = test_results
16
17 xgb.plot_importance(first_xgb)
18 plt.show()
```

Training the model..

Done. Time taken : 0:00:01.795787

Done

Evaluating the model with TRAIN data...

Evaluating Test data

TEST DATA


-----

RMSE : 1.0761851474385373

MAPE : 34.504887593204884

<IPython.core.display.Javascript object>





#### 4.4.2 Surprise BaselineModel

```
In [0]: 1 from surprise import BaselineOnly
```

##### Predicted\_rating : ( baseline prediction )

- [http://surprise.readthedocs.io/en/stable/basic\\_algorithms.html#surprise.prediction\\_algorithms.baseline\\_only.BaselineOnly](http://surprise.readthedocs.io/en/stable/basic_algorithms.html#surprise.prediction_algorithms.baseline_only.BaselineOnly)

$$\hat{r}_{ui} = b_{ui} = \mu + b_u + b_i$$

- $\mu$  : Average of all trainings in training data.
- $b_u$  : User bias
- $b_i$  : Item bias (movie biases)

##### Optimization function ( Least Squares Problem )

- [http://surprise.readthedocs.io/en/stable/prediction\\_algorithms.html#baselines-estimates-configuration](http://surprise.readthedocs.io/en/stable/prediction_algorithms.html#baselines-estimates-configuration)

$$\sum_{r_{ui} \in R_{train}} (r_{ui} - (\mu + b_u + b_i))^2 + \lambda (b_u^2 + b_i^2) . \text{ [mimimize } b_u, b_i]$$

```

In [0]: 1
2 # options are to specify.., how to compute those user and item biases
3 bsl_options = {'method': 'sgd',
4               'learning_rate': .001
5               }
6 bsl_algo = BaselineOnly(bsl_options=bsl_options)
7 # run this algorithm.., It will return the train and test results..
8 bsl_train_results, bsl_test_results = run_surprise(my_bsl_algo, trainset, tes
9
10
11 # Just store these error metrics in our models_evaluation datastructure
12 models_evaluation_train['bsl_algo'] = bsl_train_results
13 models_evaluation_test['bsl_algo'] = bsl_test_results

```

Training the model...

Estimating biases using sgd...

Done. time taken : 0:00:00.822391

Evaluating the model with train data..

time taken : 0:00:01.116752

-----

Train Data

-----

RMSE : 0.9347153928678286

MAPE : 29.389572652358183

adding train results in the dictionary..

Evaluating for test data...

time taken : 0:00:00.074418

-----

Test Data

-----

RMSE : 1.0730330260516174

MAPE : 35.04995544572911

storing the test results in test dictionary...

-----

Total time taken to run this algorithm : 0:00:02.014073

### 4.4.3 XGBoost with initial 13 features + Surprise Baseline predictor

#### Updating Train Data

```
In [0]: 1 # add our baseline_predicted value as our feature..
        2 reg_train['bslpr'] = models_evaluation_train['bsl_algo']['predictions']
        3 reg_train.head(2)
```

```
Out[44]:
```

	user	movie	GAvg	sur1	sur2	sur3	sur4	sur5	smr1	smr2	smr3	smr4	smr5	UAv
0	53406	33	3.581679	4.0	5.0	5.0	4.0	1.0	5.0	2.0	5.0	3.0	1.0	3.37037
1	99540	33	3.581679	5.0	5.0	5.0	4.0	5.0	3.0	4.0	4.0	3.0	5.0	3.55555

### Updating Test Data

```
In [0]: 1 # add that baseline predicted ratings with Surprise to the test data as well
        2 reg_test_df['bslpr'] = models_evaluation_test['bsl_algo']['predictions']
        3
        4 reg_test_df.head(2)
```

```
Out[45]:
```

	user	movie	GAvg	sur1	sur2	sur3	sur4	sur5	smr1	smr2
0	808635	71	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679
1	941866	71	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679

```
In [0]: 1 # prepare train data
2 x_train = reg_train.drop(['user', 'movie', 'rating'], axis=1)
3 y_train = reg_train['rating']
4
5 # Prepare Test data
6 x_test = reg_test_df.drop(['user', 'movie', 'rating'], axis=1)
7 y_test = reg_test_df['rating']
8
9 # initialize Our first XGBoost model...
10 xgb_bsl = xgb.XGBRegressor(silent=False, n_jobs=13, random_state=15, n_estimators=1000)
11 train_results, test_results = run_xgboost(xgb_bsl, x_train, y_train, x_test, y_test)
12
13 # store the results in models_evaluations dictionaries
14 models_evaluation_train['xgb_bsl'] = train_results
15 models_evaluation_test['xgb_bsl'] = test_results
16
17 xgb.plot_importance(xgb_bsl)
18 plt.show()
19
```

Training the model..

Done. Time taken : 0:00:02.388635

Done

Evaluating the model with TRAIN data...

Evaluating Test data


TEST DATA

-----

RMSE : 1.0763419061709816

MAPE : 34.491235560745295

<IPython.core.display.Javascript object>



#### 4.4.4 Surprise KNNBaseline predictor

In [0]: 1 `from surprise import KNNBaseline`

- KNN BASELINE

- [http://surprise.readthedocs.io/en/stable/knn\\_inspired.html#surprise.prediction\\_algorithms.knn](http://surprise.readthedocs.io/en/stable/knn_inspired.html#surprise.prediction_algorithms.knn) ([http://surprise.readthedocs.io/en/stable/knn\\_inspired.html#surprise.prediction\\_algorithms.knn](http://surprise.readthedocs.io/en/stable/knn_inspired.html#surprise.prediction_algorithms.knn))

- PEARSON\_BASELINE SIMILARITY

- [http://surprise.readthedocs.io/en/stable/similarities.html#surprise.similarities.pearson\\_baseline](http://surprise.readthedocs.io/en/stable/similarities.html#surprise.similarities.pearson_baseline) ([http://surprise.readthedocs.io/en/stable/similarities.html#surprise.similarities.pearson\\_baseline](http://surprise.readthedocs.io/en/stable/similarities.html#surprise.similarities.pearson_baseline))

- SHRINKAGE

- *2.2 Neighborhood Models* in <http://courses.ischool.berkeley.edu/i290-dm/s11/SECURE/a1-koren.pdf> (<http://courses.ischool.berkeley.edu/i290-dm/s11/SECURE/a1-koren.pdf>)

- **predicted Rating : ( \_ based on User-User similarity \_ )**

$$\hat{r}_{ui} = b_{ui} + \frac{\sum_{v \in N_i^k(u)} \text{sim}(u, v) \cdot (r_{vi} - b_{vi})}{\sum_{v \in N_i^k(u)} \text{sim}(u, v)}$$

- $b_{ui}$  - *Baseline prediction* of (user,movie) rating
- $N_i^k(u)$  - Set of **K similar** users (neighbours) of **user (u)** who rated **movie(i)**
- $\text{sim}(u, v)$  - **Similarity** between users **u** and **v**
  - Generally, it will be cosine similarity or Pearson correlation coefficient.
  - But we use **shrunk Pearson-baseline correlation coefficient**, which is based on the pearsonBaseline similarity ( we take base line predictions instead of mean rating of user/item)

- **Predicted rating** ( based on Item Item similarity ):

$$\hat{r}_{ui} = b_{ui} + \frac{\sum_{j \in N_u^k(i)} \text{sim}(i, j) \cdot (r_{uj} - b_{uj})}{\sum_{j \in N_u^k(i)} \text{sim}(i, j)}$$

- **\_Notations follows same as above (user user based predicted rating ) \_**

#### 4.4.4.1 Surprise KNNBaseline with user user similarities

```
In [0]: 1 # we specify , how to compute similarities and what to consider with sim_opti
2 sim_options = {'user_based' : True,
3               'name': 'pearson_baseline',
4               'shrinkage': 100,
5               'min_support': 2
6             }
7 # we keep other parameters like regularization parameter and learning_rate as
8 bsl_options = {'method': 'sgd'}
9
10 knn_bsl_u = KNNBaseline(k=40, sim_options = sim_options, bsl_options = bsl_op
11 knn_bsl_u_train_results, knn_bsl_u_test_results = run_surprise(knn_bsl_u, tra
12
13 # Just store these error metrics in our models_evaluation datastructure
14 models_evaluation_train['knn_bsl_u'] = knn_bsl_u_train_results
15 models_evaluation_test['knn_bsl_u'] = knn_bsl_u_test_results
16
```

Training the model...

Estimating biases using sgd...

Computing the pearson\_baseline similarity matrix...

Done computing similarity matrix.

Done. time taken : 0:00:30.173847

Evaluating the model with train data..

time taken : 0:01:35.970614

-----

Train Data

-----

RMSE : 0.33642097416508826

MAPE : 9.145093375416348

adding train results in the dictionary..

Evaluating for test data...

time taken : 0:00:00.075213

-----

Test Data

-----

RMSE : 1.0726493739667242

MAPE : 35.02094499698424

storing the test results in test dictionary...

-----

Total time taken to run this algorithm : 0:02:06.220108

#### 4.4.4.2 Surprise KNNBaseline with movie movie similarities

```

In [0]: 1 # we specify , how to compute similarities and what to consider with sim_opti
        2
        3 # 'user_based' : Fals => this considers the similarities of movies instead of
        4
        5 sim_options = {'user_based' : False,
        6               'name': 'pearson_baseline',
        7               'shrinkage': 100,
        8               'min_support': 2
        9               }
       10 # we keep other parameters like regularization parameter and learning_rate as
       11 bsl_options = {'method': 'sgd'}
       12
       13
       14 knn_bsl_m = KNNBaseline(k=40, sim_options = sim_options, bsl_options = bsl_op
       15
       16 knn_bsl_m_train_results, knn_bsl_m_test_results = run_surprise(knn_bsl_m, tra
       17
       18 # Just store these error metrics in our models_evaluation datastructure
       19 models_evaluation_train['knn_bsl_m'] = knn_bsl_m_train_results
       20 models_evaluation_test['knn_bsl_m'] = knn_bsl_m_test_results
       21

```

Training the model...

Estimating biases using sgd...

Computing the pearson\_baseline similarity matrix...

Done computing similarity matrix.

Done. time taken : 0:00:01.093096

Evaluating the model with train data..

time taken : 0:00:07.964272

-----

Train Data

-----

RMSE : 0.32584796251610554

MAPE : 8.447062581998374

adding train results in the dictionary..

Evaluating for test data...

time taken : 0:00:00.075229

-----

Test Data

-----

RMSE : 1.072758832653683

MAPE : 35.02269653015042

storing the test results in test dictionary...

-----

Total time taken to run this algorithm : 0:00:09.133017



## 4.4.5 XGBoost with initial 13 features + Surprise Baseline predictor + KNNBaseline predictor

- First we will run XGBoost with predictions from both KNN's ( that uses User\_User and Item\_Item similarities along with our previous features.
- Then we will run XGBoost with just predictions form both knn models and preditions from our baseline model.

### Preparing Train data

```
In [0]: 1 # add the predicted values from both knns to this dataframe
2 reg_train['knn_bsl_u'] = models_evaluation_train['knn_bsl_u']['predictions']
3 reg_train['knn_bsl_m'] = models_evaluation_train['knn_bsl_m']['predictions']
4
5 reg_train.head(2)
```

```
Out[51]:
```

	user	movie	GAvg	sur1	sur2	sur3	sur4	sur5	smr1	smr2	smr3	smr4	smr5	UAv
0	53406	33	3.581679	4.0	5.0	5.0	4.0	1.0	5.0	2.0	5.0	3.0	1.0	3.37037
1	99540	33	3.581679	5.0	5.0	5.0	4.0	5.0	3.0	4.0	4.0	3.0	5.0	3.55555

### Preparing Test data

```
In [0]: 1 reg_test_df['knn_bsl_u'] = models_evaluation_test['knn_bsl_u']['predictions']
2 reg_test_df['knn_bsl_m'] = models_evaluation_test['knn_bsl_m']['predictions']
3
4 reg_test_df.head(2)
```

```
Out[52]:
```

	user	movie	GAvg	sur1	sur2	sur3	sur4	sur5	smr1	smr2
0	808635	71	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679
1	941866	71	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679

```
In [0]: 1 # prepare the train data....
2 x_train = reg_train.drop(['user', 'movie', 'rating'], axis=1)
3 y_train = reg_train['rating']
4
5 # prepare the train data....
6 x_test = reg_test_df.drop(['user', 'movie', 'rating'], axis=1)
7 y_test = reg_test_df['rating']
8
9 # declare the model
10 xgb_knn_bsl = xgb.XGBRegressor(n_jobs=10, random_state=15)
11 train_results, test_results = run_xgboost(xgb_knn_bsl, x_train, y_train, x_te
12
13 # store the results in models_evaluations dictionaries
14 models_evaluation_train['xgb_knn_bsl'] = train_results
15 models_evaluation_test['xgb_knn_bsl'] = test_results
16
17
18 xgb.plot_importance(xgb_knn_bsl)
19 plt.show()
```

Training the model..

Done. Time taken : 0:00:02.092387

Done

Evaluating the model with TRAIN data...

Evaluating Test data


TEST DATA

-----

RMSE : 1.0763602465199797

MAPE : 34.48862808016984

<IPython.core.display.Javascript object>



## 4.4.6 Matrix Factorization Techniques

### 4.4.6.1 SVD Matrix Factorization User Movie interactions

```
In [0]: 1 from surprise import SVD
```

[http://surprise.readthedocs.io/en/stable/matrix\\_factorization.html#surprise.prediction\\_algorithms.matrix\\_factorization](http://surprise.readthedocs.io/en/stable/matrix_factorization.html#surprise.prediction_algorithms.matrix_factorization)  
([http://surprise.readthedocs.io/en/stable/matrix\\_factorization.html#surprise.prediction\\_algorithms.matrix\\_factorization](http://surprise.readthedocs.io/en/stable/matrix_factorization.html#surprise.prediction_algorithms.matrix_factorization))

- **Predicted Rating :**

- - $\hat{r}_{ui} = \mu + b_u + b_i + q_i^T p_u$ 
    - $q_i$  - Representation of item(movie) in latent factor space
    - $p_u$  - Representation of user in new latent factor space

- A BASIC MATRIX FACTORIZATION MODEL in [https://datajobs.com/data-science-repo/Recommender-Systems-\[Netflix\].pdf](https://datajobs.com/data-science-repo/Recommender-Systems-[Netflix].pdf) ([https://datajobs.com/data-science-repo/Recommender-Systems-\[Netflix\].pdf](https://datajobs.com/data-science-repo/Recommender-Systems-[Netflix].pdf))

- **Optimization problem with user item interactions and regularization (to avoid overfitting)**

$$\sum_{r_{ui} \in R_{train}} (r_{ui} - \hat{r}_{ui})^2 + \lambda (b_i^2 + b_u^2 + \|q_i\|^2 + \|p_u\|^2)$$

```
In [0]: 1 # initialize the model
2 from surprise import SVD
3 svd = SVD(n_factors=100, biased=True, random_state=15, verbose=True)
4 svd_train_results, svd_test_results = run_surprise(svd, trainset, testset, ve
5
6 # Just store these error metrics in our models_evaluation datastructure
7 models_evaluation_train['svd'] = svd_train_results
8 models_evaluation_test['svd'] = svd_test_results
```

Training the model...

Processing epoch 0

Processing epoch 1

Processing epoch 2

Processing epoch 3

Processing epoch 4

Processing epoch 5

Processing epoch 6

Processing epoch 7

Processing epoch 8

Processing epoch 9

Processing epoch 10

Processing epoch 11

Processing epoch 12

Processing epoch 13

Processing epoch 14

Processing epoch 15

Processing epoch 16

Processing epoch 17

Processing epoch 18

Processing epoch 19

Done. time taken : 0:00:07.297438

Evaluating the model with train data..

time taken : 0:00:01.305539

-----

Train Data

-----

RMSE : 0.6574721240954099

MAPE : 19.704901088660474

adding train results in the dictionary..

Evaluating for test data...

time taken : 0:00:00.067811

-----

Test Data

-----

RMSE : 1.0726046873826458

MAPE : 35.01953535988152

storing the test results in test dictionary...

-----

Total time taken to run this algorithm : 0:00:08.671347

#### 4.4.6.2 SVD Matrix Factorization with implicit feedback from user ( user rated movies )

```
In [0]: 1 from surprise import SVDpp
```

- ----> 2.5 Implicit Feedback in <http://courses.ischool.berkeley.edu/i290-dm/s11/SECURE/a1-koren.pdf> (<http://courses.ischool.berkeley.edu/i290-dm/s11/SECURE/a1-koren.pdf>)

- **Predicted Rating :**

- 

- $$\hat{r}_{ui} = \mu + b_u + b_i + q_i^T \left( p_u + |I_u|^{-\frac{1}{2}} \sum_{j \in I_u} y_j \right)$$

- $I_u$  --- the set of all items rated by user u
- $y_j$  --- Our new set of item factors that capture implicit ratings.
- **Optimization problem with user item interactions and regularization (to avoid overfitting)**

- 

- $$\sum_{r_{ui} \in R_{train}} (r_{ui} - \hat{r}_{ui})^2 + \lambda (b_i^2 + b_u^2 + \|q_i\|^2 + \|p_u\|^2 + \|y_j\|^2)$$

```
In [0]: 1 # initialize the model
2 svdpp = SVDpp(n_factors=50, random_state=15, verbose=True)
3 svdpp_train_results, svdpp_test_results = run_surprise(svdpp, trainset, tests
4
5 # Just store these error metrics in our models_evaluation datastructure
6 models_evaluation_train['svdpp'] = svdpp_train_results
7 models_evaluation_test['svdpp'] = svdpp_test_results
8
```

Training the model...

processing epoch 0  
processing epoch 1  
processing epoch 2  
processing epoch 3  
processing epoch 4  
processing epoch 5  
processing epoch 6  
processing epoch 7  
processing epoch 8  
processing epoch 9  
processing epoch 10  
processing epoch 11  
processing epoch 12  
processing epoch 13  
processing epoch 14  
processing epoch 15  
processing epoch 16  
processing epoch 17  
processing epoch 18  
processing epoch 19

Done. time taken : 0:01:56.765007

Evaluating the model with train data..

time taken : 0:00:06.387920

-----

Train Data

-----

RMSE : 0.6032438403305899

MAPE : 17.49285063490268

adding train results in the dictionary..

Evaluating for test data...

time taken : 0:00:00.071642

-----

Test Data

-----

RMSE : 1.0728491944183447

MAPE : 35.03817913919887

storing the test results in test dictionary...

-----

Total time taken to run this algorithm : 0:02:03.225068

## 4.4.7 XgBoost with 13 features + Surprise Baseline + Surprise KNNbaseline + MF Techniques

### Preparing Train data

```
In [0]: 1 # add the predicted values from both knns to this dataframe
        2 reg_train['svd'] = models_evaluation_train['svd']['predictions']
        3 reg_train['svdpp'] = models_evaluation_train['svdpp']['predictions']
        4
        5 reg_train.head(2)
```

```
Out[59]:
```

	user	movie	GAvg	sur1	sur2	sur3	sur4	sur5	smr1	smr2	...	smr4	smr5	UAvg
0	53406	33	3.581679	4.0	5.0	5.0	4.0	1.0	5.0	2.0	...	3.0	1.0	3.370370
1	99540	33	3.581679	5.0	5.0	5.0	4.0	5.0	3.0	4.0	...	3.0	5.0	3.555556

2 rows × 21 columns



### Preparing Test data

```
In [0]: 1 reg_test_df['svd'] = models_evaluation_test['svd']['predictions']
        2 reg_test_df['svdpp'] = models_evaluation_test['svdpp']['predictions']
        3
        4 reg_test_df.head(2)
```

```
Out[60]:
```

	user	movie	GAvg	sur1	sur2	sur3	sur4	sur5	smr1	smr2
0	808635	71	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679
1	941866	71	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679

2 rows × 21 columns





```
In [0]: 1 # prepare x_train and y_train
2 x_train = reg_train.drop(['user', 'movie', 'rating'], axis=1)
3 y_train = reg_train['rating']
4
5 # prepare test data
6 x_test = reg_test_df.drop(['user', 'movie', 'rating'], axis=1)
7 y_test = reg_test_df['rating']
8
9
10
11 xgb_final = xgb.XGBRegressor(n_jobs=10, random_state=15)
12 train_results, test_results = run_xgboost(xgb_final, x_train, y_train, x_test)
13
14 # store the results in models_evaluations dictionaries
15 models_evaluation_train['xgb_final'] = train_results
16 models_evaluation_test['xgb_final'] = test_results
17
18
19 xgb.plot_importance(xgb_final)
20 plt.show()
```

Training the model..

Done. Time taken : 0:00:04.203252

Done

Evaluating the model with TRAIN data...

Evaluating Test data


TEST DATA

-----

RMSE : 1.0763580984894978

MAPE : 34.487391651053336

<IPython.core.display.Javascript object>



#### 4.4.8 XgBoost with Surprise Baseline + Surprise KNNbaseline + MF Techniques

```
In [0]: 1 # prepare train data
2 x_train = reg_train[['knn_bsl_u', 'knn_bsl_m', 'svd', 'svdpp']]
3 y_train = reg_train['rating']
4
5 # test data
6 x_test = reg_test_df[['knn_bsl_u', 'knn_bsl_m', 'svd', 'svdpp']]
7 y_test = reg_test_df['rating']
8
9
10 xgb_all_models = xgb.XGBRegressor(n_jobs=10, random_state=15)
11 train_results, test_results = run_xgboost(xgb_all_models, x_train, y_train, x
12
13 # store the results in models_evaluations dictionaries
14 models_evaluation_train['xgb_all_models'] = train_results
15 models_evaluation_test['xgb_all_models'] = test_results
16
17 xgb.plot_importance(xgb_all_models)
18 plt.show()
```

Training the model..

Done. Time taken : 0:00:01.292225

Done

Evaluating the model with TRAIN data...

Evaluating Test data


TEST DATA

-----

RMSE : 1.075480663561971

MAPE : 35.01826709436013

<IPython.core.display.Javascript object>



## 4.5 Comparision between all models

```
In [0]: 1 # Saving our TEST_RESULTS into a dataframe so that you don't have to run it a
2 pd.DataFrame(models_evaluation_test).to_csv('sample/small/small_sample_result
3 models = pd.read_csv('sample/small/small_sample_results.csv', index_col=0)
4 models.loc['rmse'].sort_values()
```

```
Out[67]: svd                1.0726046873826458
knn_bsl_u            1.0726493739667242
knn_bsl_m            1.072758832653683
svdpp                1.0728491944183447
bsl_algo            1.0730330260516174
xgb_knn_bsl_mu      1.0753229281412784
xgb_all_models      1.075480663561971
first_algo          1.0761851474385373
xgb_bsl             1.0763419061709816
xgb_final           1.0763580984894978
xgb_knn_bsl         1.0763602465199797
Name: rmse, dtype: object
```

```
In [0]: 1 print("-"*100)
        2 print("Total time taken to run this entire notebook ( with saved files) is :"
```

```
-----
-----
Total time taken to run this entire notebook ( with saved files) is : 0:42:08.3
02761
```

## 5. Assignment

1. Instead of using 10K users and 1K movies to train the above models, use 25K users and 3K movies (or more) to train all of the above models. Report the RMSE and MAPE on the test data using larger amount of data and provide a comparison between various models as shown above.

NOTE: Please be patient as some of the code snippets make take many hours to complete execution.

2. Tune hyperparameters of all the Xgboost models above to improve the RMSE.

for the assignment part we will be considering 25k users for training data and 10k users for the test data

## getting the training and test sparse matrix

```
In [3]: 1 start = datetime.now()
        2 if os.path.isfile('train_sparse_matrix.npz'):
        3     print("It is present in your pwd, getting it from disk....")
        4     # just get it from the disk instead of computing it
        5     train_sparse_matrix = sparse.load_npz('train_sparse_matrix.npz')
        6     print("DONE..")
        7 else:
        8     print("We are creating sparse_matrix from the dataframe..")
        9     # create sparse_matrix and store it for after usage.
       10     # csr_matrix(data_values, (row_index, col_index), shape_of_matrix)
       11     # It should be in such a way that, MATRIX[row, col] = data
       12     train_sparse_matrix = sparse.csr_matrix((train_df.rating.values, (train_d
       13                                             train_df.movie.values)),)
       14
       15     print('Done. It\'s shape is : (user, movie) : ', train_sparse_matrix.shape)
       16     print('Saving it into disk for furthur usage..')
       17     # save it into disk
       18     sparse.save_npz("train_sparse_matrix.npz", train_sparse_matrix)
       19     print('Done..\n')
       20
       21 print(datetime.now() - start)
```

```
It is present in your pwd, getting it from disk....
DONE..
0:00:04.665748
```

```
In [4]: 1 start = datetime.now()
2 if os.path.isfile('test_sparse_matrix.npz'):
3     print("It is present in your pwd, getting it from disk....")
4     # just get it from the disk instead of computing it
5     test_sparse_matrix = sparse.load_npz('test_sparse_matrix.npz')
6     print("DONE..")
7 else:
8     print("We are creating sparse_matrix from the dataframe..")
9     # create sparse_matrix and store it for after usage.
10    # csr_matrix(data_values, (row_index, col_index), shape_of_matrix)
11    # It should be in such a way that, MATRIX[row, col] = data
12    test_sparse_matrix = sparse.csr_matrix((test_df.rating.values, (test_df.u
13                                           test_df.movie.values)))
14
15    print('Done. It\'s shape is : (user, movie) : ', test_sparse_matrix.shape)
16    print('Saving it into disk for further usage..')
17    # save it into disk
18    sparse.save_npz("test_sparse_matrix.npz", test_sparse_matrix)
19    print('Done..\n')
20
21 print(datetime.now() - start)
```

It is present in your pwd, getting it from disk....

DONE..

0:00:01.240002

## Sampling for trainig and test data

```

In [5]: 1 def get_sample_sparse_matrix(sparse_matrix, no_users, no_movies, path, verbose):
2         """
3         It will get it from the 'path' if it is present or It will create
4         and store the sampled sparse matrix in the path specified.
5         """
6
7         # get (row, col) and (rating) tuple from sparse_matrix...
8         row_ind, col_ind, ratings = sparse.find(sparse_matrix)
9         users = np.unique(row_ind)
10        movies = np.unique(col_ind)
11
12        print("Original Matrix : (users, movies) -- ({} {})".format(len(users), len(movies)))
13        print("Original Matrix : Ratings -- {}".format(len(ratings)))
14
15        # It just to make sure to get same sample everytime we run this program..
16        # and pick without replacement....
17        np.random.seed(15)
18        sample_users = np.random.choice(users, no_users, replace=False)
19        sample_movies = np.random.choice(movies, no_movies, replace=False)
20        # get the boolean mask or these sampled_items in originl row/col_inds..
21        mask = np.logical_and( np.isin(row_ind, sample_users),
22                               np.isin(col_ind, sample_movies) )
23
24        sample_sparse_matrix = sparse.csr_matrix((ratings[mask], (row_ind[mask],
25                                                                col_ind[mask])),
26                                                  shape=(max(sample_users)+1, max(sample_movies)+1))
27
28        if verbose:
29            print("Sampled Matrix : (users, movies) -- ({} {})".format(len(sample_users), len(sample_movies)))
30            print("Sampled Matrix : Ratings --", format(ratings[mask].shape[0]))
31
32        print('Saving it into disk for furthur usage..')
33        # save it into disk
34        sparse.save_npz(path, sample_sparse_matrix)
35        if verbose:
36            print('Done..\n')
37
38        return sample_sparse_matrix

```

**25K users and 3000 movies for train data, 13000 users and 1500 movies for test data**

```
In [6]: 1 start = datetime.now()
2 path = "sample_train_sparse_matrix.npz"
3 if os.path.isfile(path):
4     print("It is present in your pwd, getting it from disk....")
5     # just get it from the disk instead of computing it
6     sample_train_sparse_matrix = sparse.load_npz(path)
7     print("DONE..")
8 else:
9     # get 10k users and 1k movies from available data
10    sample_train_sparse_matrix = get_sample_sparse_matrix(train_sparse_matrix,
11                                                         path = path)
12
13 print(datetime.now() - start)
```

It is present in your pwd, getting it from disk....

DONE..

0:00:00.441425

```
In [7]: 1 start = datetime.now()
2
3 path = "sample_test_sparse_matrix.npz"
4 if os.path.isfile(path):
5     print("It is present in your pwd, getting it from disk....")
6     # just get it from the disk instead of computing it
7     sample_test_sparse_matrix = sparse.load_npz(path)
8     print("DONE..")
9 else:
10    # get 5k users and 500 movies from available data
11    sample_test_sparse_matrix = get_sample_sparse_matrix(test_sparse_matrix,
12                                                         path = "sample_test_sparse_m
13 print(datetime.now() - start)
```

It is present in your pwd, getting it from disk....

DONE..

0:00:00.268040

```
In [8]: 1 #for training data
2 row,col = sample_train_sparse_matrix.shape
3 elem = sample_train_sparse_matrix.count_nonzero()
4 sparsity = ((row*col) - elem)/(row*col)
5 print('the sparsity in training data is:',sparsity*100)
6
7 #for test data#for training data
8 row,col = sample_test_sparse_matrix.shape
9 elem = sample_test_sparse_matrix.count_nonzero()
10 sparsity = ((row*col) - elem)/(row*col)
11 print('the sparsity in test data is:',sparsity*100)
12
```

the sparsity in training data is: 99.9981781832153

the sparsity in test data is: 99.99984658306298

## Finding the global average of all the movie



# ratings,average user ratings,average rating per movies

## Getting global average

```
In [9]: 1 #we will store the values in a dictionary
2
3 sample_train_avg = {}
4 global_average = sample_train_sparse_matrix.sum()/sample_train_sparse_matrix.
5 sample_train_avg['globalAvg'] = global_average
6 sample_train_avg
7
```

```
Out[9]: {'globalAvg': 3.5875813607223455}
```

## Getting avergae rating per user

```
In [10]: 1 # get the user averages in dictionary (key: user_id/movie_id, value: avg rati
2
3 def get_average_ratings(sparse_matrix, of_users):
4
5     # average ratings of user/axes
6     ax = 1 if of_users else 0 # 1 - User axes, 0 - Movie axes
7
8     # ".A1" is for converting Column_Matrix to 1-D numpy array
9     sum_of_ratings = sparse_matrix.sum(axis=ax).A1
10    # Boolean matrix of ratings ( whether a user rated that movie or not)
11    is_rated = sparse_matrix!=0
12    # no of ratings that each user OR movie..
13    no_of_ratings = is_rated.sum(axis=ax).A1
14
15    # max_user and max_movie ids in sparse matrix
16    u,m = sparse_matrix.shape
17    # create a dictionary of users and their average ratings..
18    average_ratings = { i : sum_of_ratings[i]/no_of_ratings[i]
19                        for i in range(u if of_users else m)
20                        if no_of_ratings[i] !=0}
21
22    # return that dictionary of average ratings
23    return average_ratings
```

```
In [11]: 1 sample_train_avg['user'] = get_average_ratings(sample_train_sparse_matrix,of
2           print('Average rating of user 1515220 is:',sample_train_avg['user'][1515220])
```

```
Average rating of user 1515220 is: 3.92307692308
```

## Getting avergae rating per movie

```
In [12]: 1 sample_train_avg['movie'] = get_average_ratings(sample_train_sparse_matrix,
2           print('\n AVerage rating of movie 15153 :',sample_train_avg['movie'][15153])
```

Average rating of movie 15153 : 2.752

```
In [13]: 1 #number of rating in the sampled train and test data
2
3 print('ALso number of ratings in the sampled training data is:',sample_train_
4 print('number of ratings given in the sampled test data is:',sample_test_spar
```

ALso number of ratings in the sampled training data is: 856986  
number of ratings given in the sampled test data is: 72192

```
In [14]: 1 # get users, movies and ratings from our samples train sparse matrix
2 sample_train_users, sample_train_movies, sample_train_ratings = sparse.find(s
```

```
In [15]: 1 ## Reading from the trainigg data
2 reg_train = pd.read_csv('reg_train.csv', names = ['user', 'movie', 'GAvg', 's
3 reg_train.head()
```

```
Out[15]:
```

	user	movie	GAvg	sur1	sur2	sur3	sur4	sur5	smr1	smr2	smr3	smr4	smr5	UA
0	53406	33	3.581679	4.0	5.0	5.0	4.0	1.0	5.0	2.0	5.0	3.0	1.0	3.3703
1	99540	33	3.581679	5.0	5.0	5.0	4.0	5.0	3.0	4.0	4.0	3.0	5.0	3.5555
2	99865	33	3.581679	5.0	5.0	4.0	5.0	3.0	5.0	4.0	4.0	5.0	4.0	3.7142
3	101620	33	3.581679	2.0	3.0	5.0	5.0	4.0	4.0	3.0	3.0	4.0	5.0	3.5844
4	112974	33	3.581679	5.0	5.0	5.0	5.0	5.0	3.0	5.0	5.0	5.0	3.0	3.7500

## featurizing the test data

```

In [ ]: 1 start = datetime.now()
        2
        3 if os.path.isfile('sample/small/reg_test.csv'):
        4     print("It is already created...")
        5 else:
        6
        7     print('preparing {} tuples for the dataset..\\n'.format(len(sample_test_ra
        8 with open('sample/small/reg_test.csv', mode='w') as reg_data_file:
        9     count = 0
       10     for (user, movie, rating) in zip(sample_test_users, sample_test_movi
       11         st = datetime.now()
       12
       13         #----- Ratings of "movie" by similar users of "user"
       14         #print(user, movie)
       15         try:
       16             # compute the similar Users of the "user"
       17             user_sim = cosine_similarity(sample_train_sparse_matrix[user]
       18             top_sim_users = user_sim.argsort()[::-1][1:] # we are ignorin
       19             # get the ratings of most similar users for this movie
       20             top_ratings = sample_train_sparse_matrix[top_sim_users, movie
       21             # we will make it's length "5" by adding movie averages to .
       22             top_sim_users_ratings = list(top_ratings[top_ratings != 0][:5
       23             top_sim_users_ratings.extend([sample_train_averages['movie']][
       24             # print(top_sim_users_ratings, end="--")
       25
       26         except (IndexError, KeyError):
       27             # It is a new User or new Movie or there are no ratings for g
       28             ##### Cold Start Problem #####
       29             top_sim_users_ratings.extend([sample_train_averages['global']
       30             #print(top_sim_users_ratings)
       31         except:
       32             print(user, movie)
       33             # we just want KeyErrors to be resolved. Not every Exception.
       34             raise
       35
       36
       37
       38         #----- Ratings by "user" to similar movies of "m
       39         try:
       40             # compute the similar movies of the "movie"
       41             movie_sim = cosine_similarity(sample_train_sparse_matrix[:,mc
       42             top_sim_movies = movie_sim.argsort()[::-1][1:] # we are ignor
       43             # get the ratings of most similar movie rated by this user..
       44             top_ratings = sample_train_sparse_matrix[user, top_sim_movies
       45             # we will make it's length "5" by adding user averages to.
       46             top_sim_movies_ratings = list(top_ratings[top_ratings != 0][:
       47             top_sim_movies_ratings.extend([sample_train_averages['user']][
       48             #print(top_sim_movies_ratings)
       49         except (IndexError, KeyError):
       50             #print(top_sim_movies_ratings, end=" : -- ")
       51             top_sim_movies_ratings.extend([sample_train_averages['global']
       52             #print(top_sim_movies_ratings)
       53         except :
       54             raise
       55
       56         #-----prepare the row to be stores in a file-----

```

```

57     row = list()
58     # add usser and movie name first
59     row.append(user)
60     row.append(movie)
61     row.append(sample_train_averages['global']) # first feature
62     #print(row)
63     # next 5 features are similar_users "movie" ratings
64     row.extend(top_sim_users_ratings)
65     #print(row)
66     # next 5 features are "user" ratings for similar_movies
67     row.extend(top_sim_movies_ratings)
68     #print(row)
69     # Avg_user rating
70     try:
71         row.append(sample_train_averages['user'][user])
72     except KeyError:
73         row.append(sample_train_averages['global'])
74     except:
75         raise
76     #print(row)
77     # Avg_movie rating
78     try:
79         row.append(sample_train_averages['movie'][movie])
80     except KeyError:
81         row.append(sample_train_averages['global'])
82     except:
83         raise
84     #print(row)
85     # finalley, The actual Rating of this user-movie pair...
86     row.append(rating)
87     #print(row)
88     count = count + 1
89
90     # add rows to the file opened..
91     reg_data_file.write(','.join(map(str, row)))
92     #print(','.join(map(str, row)))
93     reg_data_file.write('\n')
94     if (count)%1000 == 0:
95         #print(','.join(map(str, row)))
96         print("Done for {} rows----- {}".format(count, datetime.now()))
97     print("",datetime.now() - start)

```

```
In [16]: 1 reg_test = pd.read_csv('reg_test.csv', names = ['user', 'movie', 'GAvg', 'sur1', 'sur2', 'sur3', 'sur4', 'sur5', 'smr1', 'smr2', 'rating'])
2
3
4
5 reg_test.astype({'UAvg': 'float64'}).dtypes
6 reg_test.head()
```

```
Out[16]:
```

	user	movie	GAvg	sur1	sur2	sur3	sur4	sur5	smr1	smr2
0	808635	71	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679
1	941866	71	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679
2	1737912	71	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679
3	1849204	71	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679
4	28572	111	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679	3.581679

## Transforming the data for surprise models

```
In [17]: 1 from surprise import Reader, Dataset
```

```
In [18]: 1 # It is to specify how to read the dataframe.
2 # for our dataframe, we don't have to specify anything extra..
3 start = datetime.now()#for the time factor
4 reader = Reader(rating_scale=(1,5))
5
6 # create the traindata from the dataframe...
7 train_data = Dataset.load_from_df(reg_train[['user', 'movie', 'rating']], reader)
8
9 # build the trainset from traindata.. It is of dataset format from surprise
10 trainset = train_data.build_full_trainset()
11
12 print('time taken for the computation is:', datetime.now() - start)
```

time taken for the computation is: 0:00:00.256041

## Transforming the test dataset

```
In [20]: 1 testset = list(zip(reg_test.user.values, reg_test.movie.values, reg_test.rating.values))
2 testset[:3]
```

```
Out[20]: [(808635, 71, 5), (941866, 71, 4), (1737912, 71, 3)]
```

## Applying the Machine Learning models

```
In [21]: 1 models_evaluation_train = dict()
          2 models_evaluation_test = dict()
          3
          4 models_evaluation_train, models_evaluation_test
```

```
Out[21]: ({}, {})
```

## Utility function for running the regression models

In [22]:

```

1  # to get rmse and mape given actual and predicted ratings..
2  def get_error_metrics(y_true, y_pred):
3      rmse = np.sqrt(np.mean([ (y_true[i] - y_pred[i])**2 for i in range(len(y_
4      mape = np.mean(np.abs( (y_true - y_pred)/y_true )) * 100
5      return rmse, mape
6
7  #####
8  #####
9  def run_xgboost(algo, x_train, y_train, x_test, y_test, verbose=True):
10     """
11     It will return train_results and test_results
12     """
13
14     # dictionaries for storing train and test results
15     train_results = dict()
16     test_results = dict()
17
18
19     # fit the model
20     print('Training the model..')
21     start = datetime.now()
22     algo.fit(x_train, y_train, eval_metric = 'rmse')
23     print('Done. Time taken : {}'.format(datetime.now()-start))
24     print('Done \n')
25
26     # from the trained model, get the predictions....
27     print('Evaluating the model with TRAIN data...')
28     start = datetime.now()
29     y_train_pred = algo.predict(x_train)
30     # get the rmse and mape of train data...
31     rmse_train, mape_train = get_error_metrics(y_train.values, y_train_pred)
32
33     # store the results in train_results dictionary..
34     train_results = {'rmse': rmse_train,
35                     'mape' : mape_train,
36                     'predictions' : y_train_pred}
37
38     #####
39     # get the test data predictions and compute rmse and mape
40     print('Evaluating Test data')
41     y_test_pred = algo.predict(x_test)
42     rmse_test, mape_test = get_error_metrics(y_true=y_test.values, y_pred=y_t
43     # store them in our test results dictionary.
44     test_results = {'rmse': rmse_test,
45                    'mape' : mape_test,
46                    'predictions':y_test_pred}
47
48     if verbose:
49         print('\nTEST DATA')
50         print('-'*30)
51         print('RMSE : ', rmse_test)
52         print('MAPE : ', mape_test)
53
54     # return these train and test results...
55     return train_results, test_results

```

## Utility function for running the surprise models



In [23]:

```

1  # it is just to make sure that all of our algorithms should produce same results
2  # everytime they run...
3
4  my_seed = 15
5  random.seed(my_seed)
6  np.random.seed(my_seed)
7
8  #####
9  # get (actual_list , predicted_list) ratings given list
10 # of predictions (prediction is a class in Surprise).
11 #####
12 def get_ratings(predictions):
13     actual = np.array([pred.r_ui for pred in predictions])
14     pred = np.array([pred.est for pred in predictions])
15
16     return actual, pred
17
18 #####
19 # get 'rmse' and 'mape' , given list of prediction objects
20 #####
21 def get_errors(predictions, print_them=False):
22
23     actual, pred = get_ratings(predictions)
24     rmse = np.sqrt(np.mean((pred - actual)**2))
25     mape = np.mean(np.abs(pred - actual)/actual)
26
27     return rmse, mape*100
28
29 #####
30 # It will return predicted ratings, rmse and mape of both train and test data
31 #####
32 def run_surprise(algo, trainset, testset, verbose=True):
33     '''
34         return train_dict, test_dict
35
36         It returns two dictionaries, one for train and the other is for test
37         Each of them have 3 key-value pairs, which specify 'rmse', 'mape'
38     '''
39     start = datetime.now()
40     # dictionaries that stores metrics for train and test..
41     train = dict()
42     test = dict()
43
44     # train the algorithm with the trainset
45     st = datetime.now()
46     print('Training the model...')
47     algo.fit(trainset)
48     print('Done. time taken : {} \n'.format(datetime.now()-st))
49
50     # ----- Evaluating train data-----#
51     st = datetime.now()
52     print('Evaluating the model with train data..')
53     # get the train predictions (list of prediction class inside Surprise)
54     train_preds = algo.test(trainset.build_testset())
55     # get predicted ratings from the train predictions..
56     train_actual_ratings, train_pred_ratings = get_ratings(train_preds)

```

```

57 # get 'rmse' and 'mape' from the train predictions.
58 train_rmse, train_mape = get_errors(train_preds)
59 print('time taken : {}'.format(datetime.now()-st))
60
61 if verbose:
62     print('-'*15)
63     print('Train Data')
64     print('-'*15)
65     print("RMSE : {}\nMAPE : {}".format(train_rmse, train_mape))
66
67 #store them in the train dictionary
68 if verbose:
69     print('adding train results in the dictionary..')
70 train['rmse'] = train_rmse
71 train['mape'] = train_mape
72 train['predictions'] = train_pred_ratings
73
74 #----- Evaluating Test data-----#
75 st = datetime.now()
76 print('\nEvaluating for test data...')
77 # get the predictions( list of prediction classes) of test data
78 test_preds = algo.test(testset)
79 # get the predicted ratings from the list of predictions
80 test_actual_ratings, test_pred_ratings = get_ratings(test_preds)
81 # get error metrics from the predicted and actual ratings
82 test_rmse, test_mape = get_errors(test_preds)
83 print('time taken : {}'.format(datetime.now()-st))
84
85 if verbose:
86     print('-'*15)
87     print('Test Data')
88     print('-'*15)
89     print("RMSE : {}\nMAPE : {}".format(test_rmse, test_mape))
90 # store them in test dictionary
91 if verbose:
92     print('storing the test results in test dictionary...')
93 test['rmse'] = test_rmse
94 test['mape'] = test_mape
95 test['predictions'] = test_pred_ratings
96
97 print('\n'+ '-'*45)
98 print('Total time taken to run this algorithm :', datetime.now() - start)
99
100 # return two dictionaries train and test
101 return train, test

```

## XGBOOST with 13 features

In [24]:

```

1 import xgboost as XGBRegressor
2 from sklearn.model_selection import GridSearchCV
3 from sklearn.model_selection import TimeSeriesSplit

```

In [25]:

```

1 import warnings
2 warnings.filterwarnings('ignore')

```

```
In [26]: 1 x_train = reg_train.drop(['user', 'movie', 'rating'], axis = 1)
          2 y_train = reg_train['rating']
          3
          4 #for test data
          5 x_test = reg_test.drop(['user', 'movie', 'rating'], axis = 1)
          6 y_test = reg_test['rating']
```

```
In [34]: 1 import xgboost as xgb
2 params = {}
3 #params['objective'] = 'reg:squarederror'
4 params['eval_metric'] = 'rmse'
5 params['eta'] = 0.02
6 params['max_depth'] = 3
7 params['colsample_bytree'] = 0.7
8 params['n_estimators'] = 1100
9 params['subsample'] = 0.3
10 params['learning_rate'] = 0.1
11
12 d_train = xgb.DMatrix(x_train, label=y_train)
13 d_test = xgb.DMatrix(x_test, label = y_test)
14
15 watchlist = [(d_train, 'train'), (d_test, 'valid')]
16
17 bst = xgb.train(params, d_train, 400, watchlist, verbose_eval= 10, early_stoppi
18
19 xgdmatrix = xgb.DMatrix(x_train, y_train)
20 predict_train = bst.predict(d_train)
21 predict_test = bst.predict(d_test)
22
23 rmse_train, mape_train = get_error_metrics(y_train, predict_train)
24 rmse_test, mape_test = get_error_metrics(y_test, predict_test)
25 print('\nThe Training rmse is :', rmse_train)
26 print('Training mape is:', mape_train)
27 print("\nThe Test rmse is :", rmse_test)
28 print('The Test mape is:', mape_test)
29
```

```
[0]    train-rmse:2.96697    valid-rmse:2.98091
```

Multiple eval metrics have been passed: 'valid-rmse' will be used for early stopping.

Will train until valid-rmse hasn't improved in 20 rounds.


```
[10]    train-rmse:1.33474    valid-rmse:1.49258
[20]    train-rmse:0.940171   valid-rmse:1.1453
[30]    train-rmse:0.869937   valid-rmse:1.08875
[40]    train-rmse:0.855913   valid-rmse:1.07781
[50]    train-rmse:0.851634   valid-rmse:1.07586
[60]    train-rmse:0.84987    valid-rmse:1.07524
[70]    train-rmse:0.848754   valid-rmse:1.07489
[80]    train-rmse:0.848063   valid-rmse:1.0748
[90]    train-rmse:0.847364   valid-rmse:1.07461
[100]   train-rmse:0.846911   valid-rmse:1.0744
[110]   train-rmse:0.846477   valid-rmse:1.07479
Stopping. Best iteration:
[96]    train-rmse:0.847077    valid-rmse:1.07437
```

```
The Training rmse is : 0.846216109815
Training mape is: 25.22057592868805
```

```
The Test rmse is : 1.07474464639
The Test mape is: 34.6695352057443
```

```
In [40]: 1 import matplotlib.pyplot as plt
2 xgb.plot_importance(bst)
3 plt.show()
```

<IPython.core.display.Javascript object>



**tuning the hyperparameters using randomizedsearchcv**

```

In [41]: 1 from xgboost import XGBRegressor
2 #Declaring parameters
3 params = {'learning_rate':[0.1,0.01,0.001,0.0001],
4           'n_estimators':[250,500,700,750,1000,1500,2000,3000],
5           'subsample':[0.6,0.7,0.8,0.9],
6           'min_child_weight':[3,5,7,9],
7           'reg_lambda':[100,200,300,400],
8           'reg_alpha':[100,200,300, 400],
9           'max_depth': [1,3,4,5,6,7,9],
10          'colsample_bytree':[0.6,0.7,0.8],
11          'gamma':[0,0.5,1]}
12
13 #Tuning hyperparameters
14 start =datetime.now()
15 model= XGBRegressor(random_state=0,n_jobs=-1,objective = 'reg:squarederror')
16 rsearch = RandomizedSearchCV(model,params,n_iter=20,scoring='neg_mean_absolut
17 rsearch.fit(x_train, y_train)
18 print('time taken to perform Hyperparameter tunings :',datetime.now()-start)
19
20 #Getting the best hyperparameter tuned model
21 best_model=rsearch.best_estimator_
22 print("Best estimator: ",best_model)
23
24 #Fitting the best model to our training data
25 #best_model.fit(df_train, tsne_train_output)
26
27

```

time taken to perform hyperparameter tunings is: 18:44:02.229826

The best estimator is : XGBRegressor(base\_score=0.5, booster='gbtree', colsample\_bylevel=1,

```

    colsample_bynode=1, colsample_bytree=1, gamma=0,
    importance_type='gain', learning_rate=0.1, max_delta_step=0,
    max_depth=3, min_child_weight=1, missing=None, n_estimators=700,
    n_jobs=1, nthread=-1, objective='reg:squarederror', random_state=0,
    reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,
    silent=None, subsample=1, verbosity=1)

```

```
In [43]: 1 train_results, test_results = run_xgboost(best_model, x_train, y_train, x_test)
2
3 # store the results in models_evaluations dictionaries
4 models_evaluation_train['first_algo'] = train_results
5 models_evaluation_test['first_algo'] = test_results
6
7 xgb.plot_importance(best_model)
8 plt.show()
```

Training the model..

Done. Time taken : 0:00:37.424939

Done

Evaluating the model with TRAIN data...

Evaluating Test data


TEST DATA

-----

RMSE : 1.07899932598

MAPE : 34.3041145148

<IPython.core.display.Javascript object>



## Surprise baseline model

```
In [45]: 1 from surprise import BaselineOnly #importing the important library
```



```
In [46]: 1 # options are to specify.., how to compute those user and item biases
2 bsl_options = {'method': 'sgd',
3               'learning_rate': .001
4               }
5 bsl_algo = BaselineOnly(bsl_options=bsl_options)
6 # run this algorithm.., It will return the train and test results..
7 bsl_train_results, bsl_test_results = run_surprise(bsl_algo, trainset, testse
8
9
10 # Just store these error metrics in our models_evaluation datastructure
11 models_evaluation_train['bsl_algo'] = bsl_train_results
12 models_evaluation_test['bsl_algo'] = bsl_test_results
```

Training the model...

Estimating biases using sgd...

Done. time taken : 0:00:00.760967

Evaluating the model with train data..

time taken : 0:00:00.965418

-----

Train Data

-----

RMSE : 0.9347153928678286

MAPE : 29.389572652358183

adding train results in the dictionary..

Evaluating for test data...

time taken : 0:00:00.119681

-----

Test Data

-----

RMSE : 1.0730330260516174

MAPE : 35.04995544572911

storing the test results in test dictionary...

-----

Total time taken to run this algorithm : 0:00:01.847063

## XGBoost with 13 initial features and baseline surprise feature

```
In [53]: 1 models_evaluation_train.keys()
```

```
Out[53]: dict_keys(['first_algo', 'bsl_algo'])
```

```
In [54]: 1 x_train['bslpr'] = models_evaluation_train['bsl_algo']['predictions']
2 x_test['bslpr'] = models_evaluation_test['bsl_algo']['predictions']
```

```

In [55]: 1 #Declaring parameters
2 params = {'learning_rate':[0.1,0.01,0.001,0.0001],
3           'n_estimators':[250,500,700,750,1000,1500,2000,3000],
4           'subsample':[0.6,0.7,0.8,0.9],
5           'min_child_weight':[3,5,7,9],
6           'reg_lambda':[100,200,300,400],
7           'reg_alpha':[100,200,300, 400],
8           'max_depth': [1,3,4,5,6,7,9],
9           'colsample_bytree':[0.6,0.7,0.8],
10          'gamma':[0,0.5,1]}
11
12 #Tuning hyperparameters
13
14 #print('Hyperparameter tuning: \n')
15 model= XGBRegressor(random_state=0,n_jobs=-1,objective = 'reg:squarederror')
16 rsearch = RandomizedSearchCV(model,params,n_iter=20,scoring='neg_mean_absolut
17 rsearch.fit(x_train, y_train)
18 #print('time taken to perform Hyperparameter tunings :',datetime.now()-start)
19
20 #Getting the best hyperparameter tuned model
21 best_model_bsl=rsearch.best_estimator_
22 print("best estimator is: ",best_model_bsl)
23

```

```

best estimator is: XGBRegressor(base_score=0.5, booster='gbtree', colsample_byl
evel=1,
    colsample_bynode=1, colsample_bytree=1, gamma=0,
    importance_type='gain', learning_rate=0.1, max_delta_step=0,
    max_depth=3, min_child_weight=1, missing=None, n_estimators=500,
    n_jobs=1, nthread=-1, objective='reg:squarederror', random_state=0,
    reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,
    silent=None, subsample=1, verbosity=1)

```

```
In [57]: 1 train_results, test_results = run_xgboost(xgb_bsl, x_train, y_train, x_test,
2
3 # store the results in models_evaluations dictionaries
4 models_evaluation_train['xgb_baseline'] = train_results
5 models_evaluation_test['xgb_baseline'] = test_results
6
7 xgb.plot_importance(xgb_bsl)
8 plt.show()
```

Training the model..

Done. Time taken : 0:00:33.068586

Done

Evaluating the model with TRAIN data...


Evaluating Test data

TEST DATA

-----  
RMSE : 1.07556227918

MAPE : 34.5682904007

<IPython.core.display.Javascript object>



## Surprise KNN Baseline predictor

```
In [58]: 1 from surprise import KNNBaseline
```

```
In [59]: 1 # we specify , how to compute similarities and what to consider with sim_opti
2 sim_options = {'user_based' : True,
3               'name': 'pearson_baseline',
4               'shrinkage': 100,
5               'min_support': 2
6             }
7 # we keep other parameters like regularization parameter and learning_rate as
8 bsl_options = {'method': 'sgd'}
9
10 knn_bsl_u = KNNBaseline(k=40, sim_options = sim_options, bsl_options = bsl_op
11 knn_bsl_train_results, knn_bsl_test_results = run_surprise(knn_bsl_u, trainse
12
13 # Just store these error metrics in our models_evaluation datastructure
14 models_evaluation_train['knn_bsl'] = knn_bsl_train_results
15 models_evaluation_test['knn_bsl'] = knn_bsl_test_results
```

Training the model...

Estimating biases using sgd...

Computing the pearson\_baseline similarity matrix...

Done computing similarity matrix.

Done. time taken : 0:00:45.483651

Evaluating the model with train data..

time taken : 0:01:39.043199

-----

Train Data

-----

RMSE : 0.33642097416508826

MAPE : 9.145093375416348

adding train results in the dictionary..

Evaluating for test data...

time taken : 0:00:00.115729

-----

Test Data

-----

RMSE : 1.0726493739667242

MAPE : 35.02094499698424

storing the test results in test dictionary...

-----

Total time taken to run this algorithm : 0:02:24.642579

## Surprise KNN Baseline for movie movie similarity

```

In [60]: 1 # we specify , how to compute similarities and what to consider with sim_opti
2
3 # 'user_based' : Fals => this considers the similarities of movies instead of
4
5 sim_options = {'user_based' : False,
6               'name': 'pearson_baseline',
7               'shrinkage': 100,
8               'min_support': 2
9             }
10 # we keep other parameters like regularization parameter and learning_rate as
11 bsl_options = {'method': 'sgd'}
12
13
14 knn_bsl_m = KNNBaseline(k=40, sim_options = sim_options, bsl_options = bsl_op
15
16 knn_bsl_m_train_results, knn_bsl_m_test_results = run_surprise(knn_bsl_m, tra
17
18 # Just store these error metrics in our models_evaluation datastructure
19 models_evaluation_train['knn_bsl_m'] = knn_bsl_m_train_results
20 models_evaluation_test['knn_bsl_m'] = knn_bsl_m_test_results
21

```

Training the model...

Estimating biases using sgd...

Computing the pearson\_baseline similarity matrix...

Done computing similarity matrix.

Done. time taken : 0:00:01.699455

Evaluating the model with train data..

time taken : 0:00:08.877304

-----

Train Data

-----

RMSE : 0.32584796251610554

MAPE : 8.447062581998374

adding train results in the dictionary..

Evaluating for test data...

time taken : 0:00:00.170747

-----

Test Data

-----

RMSE : 1.072758832653683

MAPE : 35.02269653015042

storing the test results in test dictionary...

-----

Total time taken to run this algorithm : 0:00:10.748503

## XGboost with 13 features + Surprise Baseline + Surprise

## KNN Baseline features

so first we will train the xgboost model with both features of users and movies along with 13 features, then we will train the xgboost model with 13 features + 2 Surprise KNN features + Surprise baseline features

```
In [61]: 1 x_train['knn_bsl_u'] = models_evaluation_train['knn_bsl']['predictions']
          2 x_test['knn_bsl_u'] = models_evaluation_test['knn_bsl']['predictions']
          3 x_train['knn_bsl_m'] = models_evaluation_train['knn_bsl_m']['predictions']
          4 x_test['knn_bsl_m'] = models_evaluation_test['knn_bsl_m']['predictions']
```

```
In [62]: 1 x_train.head(2)
```

```
Out[62]:
```

	GAvg	sur1	sur2	sur3	sur4	sur5	smr1	smr2	smr3	smr4	smr5	UAvg	MAvg	
0	3.581679	4.0	5.0	5.0	4.0	1.0	5.0	2.0	5.0	3.0	1.0	3.370370	4.092437	3.
1	3.581679	5.0	5.0	5.0	4.0	5.0	3.0	4.0	4.0	3.0	5.0	3.555556	4.092437	3.

## Tuning the models

```

In [64]: 1 #Declaring parameters
2 params = {'learning_rate':[0.1,0.01,0.001,0.0001],
3           'n_estimators':[250,500,700,750,1000,1500,2000,3000],
4           'subsample':[0.6,0.7,0.8,0.9],
5           'min_child_weight':[3,5,7,9],
6           'reg_lambda':[100,200,300,400],
7           'reg_alpha':[100,200,300, 400],
8           'max_depth': [1,3,4,5,6,7,9],
9           'colsample_bytree':[0.6,0.7,0.8],
10          'gamma':[0,0.5,1]}
11
12 #Tuning hyperparameters
13
14 #print('Hyperparameter tuning: \n')
15 model= XGBRegressor(random_state=0,n_jobs=-1,objective = 'reg:squarederror')
16 rsearch = RandomizedSearchCV(model,params,n_iter=20,scoring='neg_mean_absolut
17 rsearch.fit(x_train, y_train)
18 #print('time taken to perform Hyperparameter tunings :',datetime.now()-start)
19
20 #Getting the best hyperparameter tuned model
21 xgb_with_knn=rsearch.best_estimator_
22 print("best estimator is: ",xgb_with_knn)
23

```

```

best estimator is: XGBRegressor(base_score=0.5, booster='gbtree', colsample_byl
evel=1,
    colsample_bynode=1, colsample_bytree=1, gamma=0,
    importance_type='gain', learning_rate=0.1, max_delta_step=0,
    max_depth=3, min_child_weight=1, missing=None, n_estimators=300,
    n_jobs=1, nthread=-1, objective='reg:squarederror', random_state=0,
    reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,
    silent=None, subsample=1, verbosity=1)

```



```
In [66]: 1 train_results, test_results = run_xgboost(xgb_with_knn, x_train, y_train, x_t
2
3 # store the results in models_evaluations dictionaries
4 models_evaluation_train['xgb_knn_bsl'] = train_results
5 models_evaluation_test['xgb_knn_bsl'] = test_results
6
7
8 xgb.plot_importance(xgb_with_knn)
9 plt.show()
```

Training the model..

Done. Time taken : 0:00:18.476601

Done

Evaluating the model with TRAIN data...

Evaluating Test data


TEST DATA

-----

RMSE : 1.07650440946

MAPE : 34.471927498

<IPython.core.display.Javascript object>



## **Matrix factorization techniques**

### **SVD Matrix factorization for user movie interactions**

```
In [67]: 1 # initialize the model
2 from surprise import SVD
3 svd = SVD(n_factors=100, biased=True, random_state=15, verbose=True)
4 svd_train_results, svd_test_results = run_surprise(svd, trainset, testset, ve
5
6 # Just store these error metrics in our models_evaluation datastructure
7 models_evaluation_train['svd'] = svd_train_results
8 models_evaluation_test['svd'] = svd_test_results
```

Training the model...

Processing epoch 0

Processing epoch 1

Processing epoch 2

Processing epoch 3

Processing epoch 4

Processing epoch 5

Processing epoch 6

Processing epoch 7

Processing epoch 8

Processing epoch 9

Processing epoch 10

Processing epoch 11

Processing epoch 12

Processing epoch 13

Processing epoch 14

Processing epoch 15

Processing epoch 16

Processing epoch 17

Processing epoch 18

Processing epoch 19

Done. time taken : 0:00:09.700443

Evaluating the model with train data..

time taken : 0:00:01.550846

-----

Train Data

-----

RMSE : 0.6574721240954099

MAPE : 19.704901088660478

adding train results in the dictionary..

Evaluating for test data...

time taken : 0:00:00.255315

-----

Test Data

-----

RMSE : 1.0726046873826458

MAPE : 35.01953535988152

storing the test results in test dictionary...

-----

Total time taken to run this algorithm : 0:00:11.506604

## SVD matrix factorization with implicit feedback

```
In [68]: 1 from surprise import SVDpp
2         # initialize the model
3         svdpp = SVDpp(n_factors=50, random_state=15, verbose=True)
4         svdpp_train_results, svdpp_test_results = run_surprise(svdpp, trainset, tests
5
6         # Just store these error metrics in our models_evaluation datastructure
7         models_evaluation_train['svdpp'] = svdpp_train_results
8         models_evaluation_test['svdpp'] = svdpp_test_results
9
```

Training the model...

processing epoch 0  
processing epoch 1  
processing epoch 2  
processing epoch 3  
processing epoch 4  
processing epoch 5  
processing epoch 6  
processing epoch 7  
processing epoch 8  
processing epoch 9  
processing epoch 10  
processing epoch 11  
processing epoch 12  
processing epoch 13  
processing epoch 14  
processing epoch 15  
processing epoch 16  
processing epoch 17  
processing epoch 18  
processing epoch 19

Done. time taken : 0:03:07.454325

Evaluating the model with train data..

time taken : 0:00:08.546879

-----

Train Data

-----

RMSE : 0.6032438403305899

MAPE : 17.49285063490268

adding train results in the dictionary..

Evaluating for test data...

time taken : 0:00:00.151537

-----

Test Data

-----

RMSE : 1.0728491944183447

MAPE : 35.03817913919887

storing the test results in test dictionary...

-----  
Total time taken to run this algorithm : 0:03:16.154738

#### 4.4.7 XgBoost with 13 features + Surprise Baseline + Surprise KNNbaseline + MF Techniques

```
In [70]: 1 x_train['svd'] = models_evaluation_train['svd']['predictions']
2 x_train['svd_pp'] = models_evaluation_train['svdpp']['predictions']
3
4 #for test data
5 x_test['svd'] = models_evaluation_test['svd']['predictions']
6 x_test['svd_pp'] = models_evaluation_test['svdpp']['predictions']
```

```
In [71]: 1 #Declaring parameters
2 params = {'learning_rate':[0.1,0.01,0.001,0.0001],
3          'n_estimators':[250,500,700,750,1000,1500,2000,3000],
4          'subsample':[0.6,0.7,0.8,0.9],
5          'min_child_weight':[3,5,7,9],
6          'reg_lambda':[100,200,300,400],
7          'reg_alpha':[100,200,300, 400],
8          'max_depth': [1,3,4,5,6,7,9],
9          'colsample_bytree':[0.6,0.7,0.8],
10         'gamma':[0,0.5,1]}
11
12 #Tuning hyperparameters
13
14 #print('Hyperparameter tuning: \n')
15 model= XGBRegressor(random_state=0,n_jobs=-1,objective = 'reg:squarederror')
16 rsearch = RandomizedSearchCV(model,params,n_iter=20,scoring='neg_mean_absolut
17 rsearch.fit(x_train, y_train)
18 #print('time taken to perform Hyperparameter tunings :',datetime.now()-start)
19
20 #Getting the best hyperparameter tuned model
21 xgb_with_all=rsearch.best_estimator_
22 print("best estimator is: ",xgb_with_all)
```

```
best estimator is: XGBRegressor(base_score=0.5, booster='gbtree', colsample_byl
evel=1,
    colsample_bynode=1, colsample_bytree=1, gamma=0,
    importance_type='gain', learning_rate=0.1, max_delta_step=0,
    max_depth=3, min_child_weight=1, missing=None, n_estimators=500,
    n_jobs=1, nthread=-1, objective='reg:squarederror', random_state=0,
    reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,
    silent=None, subsample=1, verbosity=1)
```

```
In [72]: 1 train_results, test_results = run_xgboost(xgb_with_all, x_train, y_train, x_t
2
3 # store the results in models_evaluations dictionaries
4 models_evaluation_train['xgb_final'] = train_results
5 models_evaluation_test['xgb_final'] = test_results
6
7 xgb.plot_importance(xgb_with_all)
8 plt.show()
```

Training the model..

Done. Time taken : 0:00:52.825762

Done

Evaluating the model with TRAIN data...

Evaluating Test data


TEST DATA

-----

RMSE : 1.07623153253

MAPE : 34.5090296755

<IPython.core.display.Javascript object>



### 4.4.8 XgBoost with Surprise Baseline + Surprise KNNbaseline + MF Techniques

In [73]:

```
1 x_train.head()
```

Out[73]:

	GAvg	sur1	sur2	sur3	sur4	sur5	smr1	smr2	smr3	smr4	smr5	UAvg	MAvg	
0	3.581679	4.0	5.0	5.0	4.0	1.0	5.0	2.0	5.0	3.0	1.0	3.370370	4.092437	3.
1	3.581679	5.0	5.0	5.0	4.0	5.0	3.0	4.0	4.0	3.0	5.0	3.555556	4.092437	3.
2	3.581679	5.0	5.0	4.0	5.0	3.0	5.0	4.0	4.0	5.0	4.0	3.714286	4.092437	3.
3	3.581679	2.0	3.0	5.0	5.0	4.0	4.0	3.0	3.0	4.0	5.0	3.584416	4.092437	3.
4	3.581679	5.0	5.0	5.0	5.0	5.0	3.0	5.0	5.0	5.0	3.0	3.750000	4.092437	3.

In [74]:

```
1 X_final_train = x_train[['knn_bsl_u', 'knn_bsl_m', 'svd', 'svd_pp']]
2 X_final_test = x_test[['knn_bsl_u', 'knn_bsl_m', 'svd', 'svd_pp']]
```



```

In [75]: 1 #Declaring parameters
2 params = {'learning_rate':[0.1,0.01,0.001,0.0001],
3           'n_estimators':[250,500,700,750,1000,1500,2000,3000],
4           'subsample':[0.6,0.7,0.8,0.9],
5           'min_child_weight':[3,5,7,9],
6           'reg_lambda':[100,200,300,400],
7           'reg_alpha':[100,200,300, 400],
8           'max_depth': [1,3,4,5,6,7,9],
9           'colsample_bytree':[0.6,0.7,0.8],
10          'gamma':[0,0.5,1]}
11
12 #Tuning hyperparameters
13
14 #print('Hyperparameter tuning: \n')
15 model= XGBRegressor(random_state=0,n_jobs=-1,objective = 'reg:squarederror')
16 rsearch = RandomizedSearchCV(model,params,n_iter=20,scoring='neg_mean_absolut
17 rsearch.fit(X_final_train, y_train)
18 #print('time taken to perform Hyperparameter tunings :',datetime.now()-start)
19
20 #Getting the best hyperparameter tuned model
21 xgb_final=rsearch.best_estimator_
22 print("best estimator is: ",xgb_final)

```

```

best estimator is: XGBRegressor(base_score=0.5, booster='gbtree', colsample_byl
evel=1,
    colsample_bynode=1, colsample_bytree=1, gamma=0,
    importance_type='gain', learning_rate=0.1, max_delta_step=0,
    max_depth=1, min_child_weight=1, missing=None, n_estimators=100,
    n_jobs=1, nthread=-1, objective='reg:squarederror', random_state=0,
    reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,
    silent=None, subsample=1, verbosity=1)

```

```
In [76]: 1 train_results, test_results = run_xgboost(xgb_final,X_final_train, y_train,X_
2
3 # store the results in models_evaluations dictionaries
4 models_evaluation_train['xgb_all_models'] = train_results
5 models_evaluation_test['xgb_all_models'] = test_results
6
7 xgb.plot_importance(xgb_final)
8 plt.show()
```

Training the model..

Done. Time taken : 0:00:03.016979

Done

Evaluating the model with TRAIN data...

Evaluating Test data


TEST DATA

-----

RMSE : 1.07517647094

MAPE : 35.1258123364

<IPython.core.display.Javascript object>



## final comparison between all the models

```
In [97]: 1 # Saving our TEST_RESULTS into a dataframe so that you don't have to run it a
2 pd.DataFrame(models_evaluation_test).to_csv('final_results.csv')
3 models = pd.read_csv('final_results.csv', index_col=0)
4 print('root mean squared error is:')
5 print(models.loc['rmse'].sort_values())
6
```

```
root mean squared error is:
svd                1.07260468738
knn_bsl            1.07264937397
knn_bsl_m          1.07275883265
svdpp              1.07284919442
bsl_algo           1.07303302605
xgb_all_models     1.07517647094
xgb_baseline       1.07556227918
xgb_final          1.07623153253
xgb_knn_bsl        1.07650440946
first_algo         1.07899932598
Name: rmse, dtype: object
```

## Conclusion

In [96]:

```

1 from prettytable import PrettyTable
2 table = PrettyTable()
3 table.field_names = ['Model', 'Rmse', 'Mape']
4 table.add_row(['svd',models.loc['rmse']['svd'],models.loc['mape']['svd']])
5 table.add_row(['knn_bsl',models.loc['rmse']['knn_bsl'],models.loc['mape']['knn_bsl']])
6 table.add_row(['knn_bsl_m',models.loc['rmse']['knn_bsl_m'],models.loc['mape']['knn_bsl_m']])
7 table.add_row(['svdpp',models.loc['rmse']['svdpp'],models.loc['mape']['svdpp']])
8 table.add_row(['bsl_algo',models.loc['rmse']['bsl_algo'],models.loc['mape']['bsl_algo']])
9 table.add_row(['xgb_all_models',models.loc['rmse']['xgb_all_models'],models.loc['mape']['xgb_all_models']])
10 table.add_row(['xgb_baseline',models.loc['rmse']['xgb_baseline'],models.loc['mape']['xgb_baseline']])
11 table.add_row(['xgb_final',models.loc['rmse']['xgb_final'],models.loc['mape']['xgb_final']])
12 table.add_row(['xgb_knn_bsl',models.loc['rmse']['xgb_knn_bsl'],models.loc['mape']['xgb_knn_bsl']])
13 table.add_row(['first_algo',models.loc['rmse']['first_algo'],models.loc['mape']['first_algo']])
14 print('\tAll the models that we implemented')
15 print(table)

```

All the models that we implemented

Model	Rmse	Mape
svd	1.07260468738	35.0195353599
knn_bsl	1.07264937397	35.020944997
knn_bsl_m	1.07275883265	35.0226965302
svdpp	1.07284919442	35.0381791392
bsl_algo	1.07303302605	35.0499554457
xgb_all_models	1.07517647094	35.1258123364
xgb_baseline	1.07556227918	34.5682904007
xgb_final	1.07623153253	34.5090296755
xgb_knn_bsl	1.07650440946	34.471927498
first_algo	1.07899932598	34.3041145148

## Case study and our approach

In this case study, the business problem we were trying to solve is how to improve the netflix algorithm for recommending movies to the users, so the analysis was done in keeping in mind that how netflix has laid down certain norms and the winning solution of the team lead by professor 'Yehuda koren'.

So we approached this problem both as a regression problem and a recommendation problem by primarily converting for the sparse matrix where every user has given some or the other rating to a movie and not given rating to most of the movies, which we are trying to find through our approach and thus recommending movies which user is likely to give maximum rating thus making it a matrix completing (recommender systems) and a Regression problem (reducing the mean absolute percentage error).

We considered 25000 users and 3000 movies in training data while 13000 users and 1500 movies in the test data, also one of the important factors to keep in mind while approaching was the 'cold start problem' where we do not have any knowledge about new user and his preferences where we then compute all its similarities with the training data and then retrain whole model for next incoming new user or other different strategies.

## Featurization technique

Each row in the train dataframe will consist of a user, the movie he/she has rated, the global average of all ratings given by all the 25K users, it will also contain the ratings of top 5 similar users who has rated the movie (sur1,sur2...sur5). It has the ratings of the top 5 most similar movies to the given movie(smr1,smr2...smr5). Each row will also contain the user's average rating on all the movies he/she has watched, the average rating for the given movie and lastly the rating given by the query user on this movie.

## Overview of our modelling strategy

For recommendation systems there is an extremely fast and scalable library that we will use in order to build our models. At first, we have posed this movie recommendation problem as a regression problem. Then we use the surprise library to create a baseline model, we will use the output of this as a feature to our regression model.

We have 13 handcrafted features. We have 1 feature from the output of the surprise baseline model. We have one feature from the output of the baseline KNN model. We have another feature from the KNN movie movie similarity. We have two features from the outputs of the baseline SVD and SVD++ models. We will use all these outputs as features to the regression model we will build. One important note we have to keep in mind is that we cannot use the test data for feature engineering. Suppose there's a new user who has subscribed to Netflix. We don't have prior data about the user, so it's a cold start problem. This means we cannot use his/her data for feature engineering. In case of a new user we will put the value of engineered features to be zero. Logically it is the right thing to do.

1. performed XGboost with 13 features
2. Then on XGBoost with initial 13 features + Surprise Baseline predictor.
3. Then on XGBoost with initial 13 features + Surprise Baseline predictor + KNNBaseline predictor.
4. Also XGBoost with initial 13 features + Surprise Baseline predictor + KNNBaseline predictor + SVD.
5. Also XGBoost with initial 13 features , SVD ,SVD++, Surprise Baseline predictor + KNNBaseline predictor.